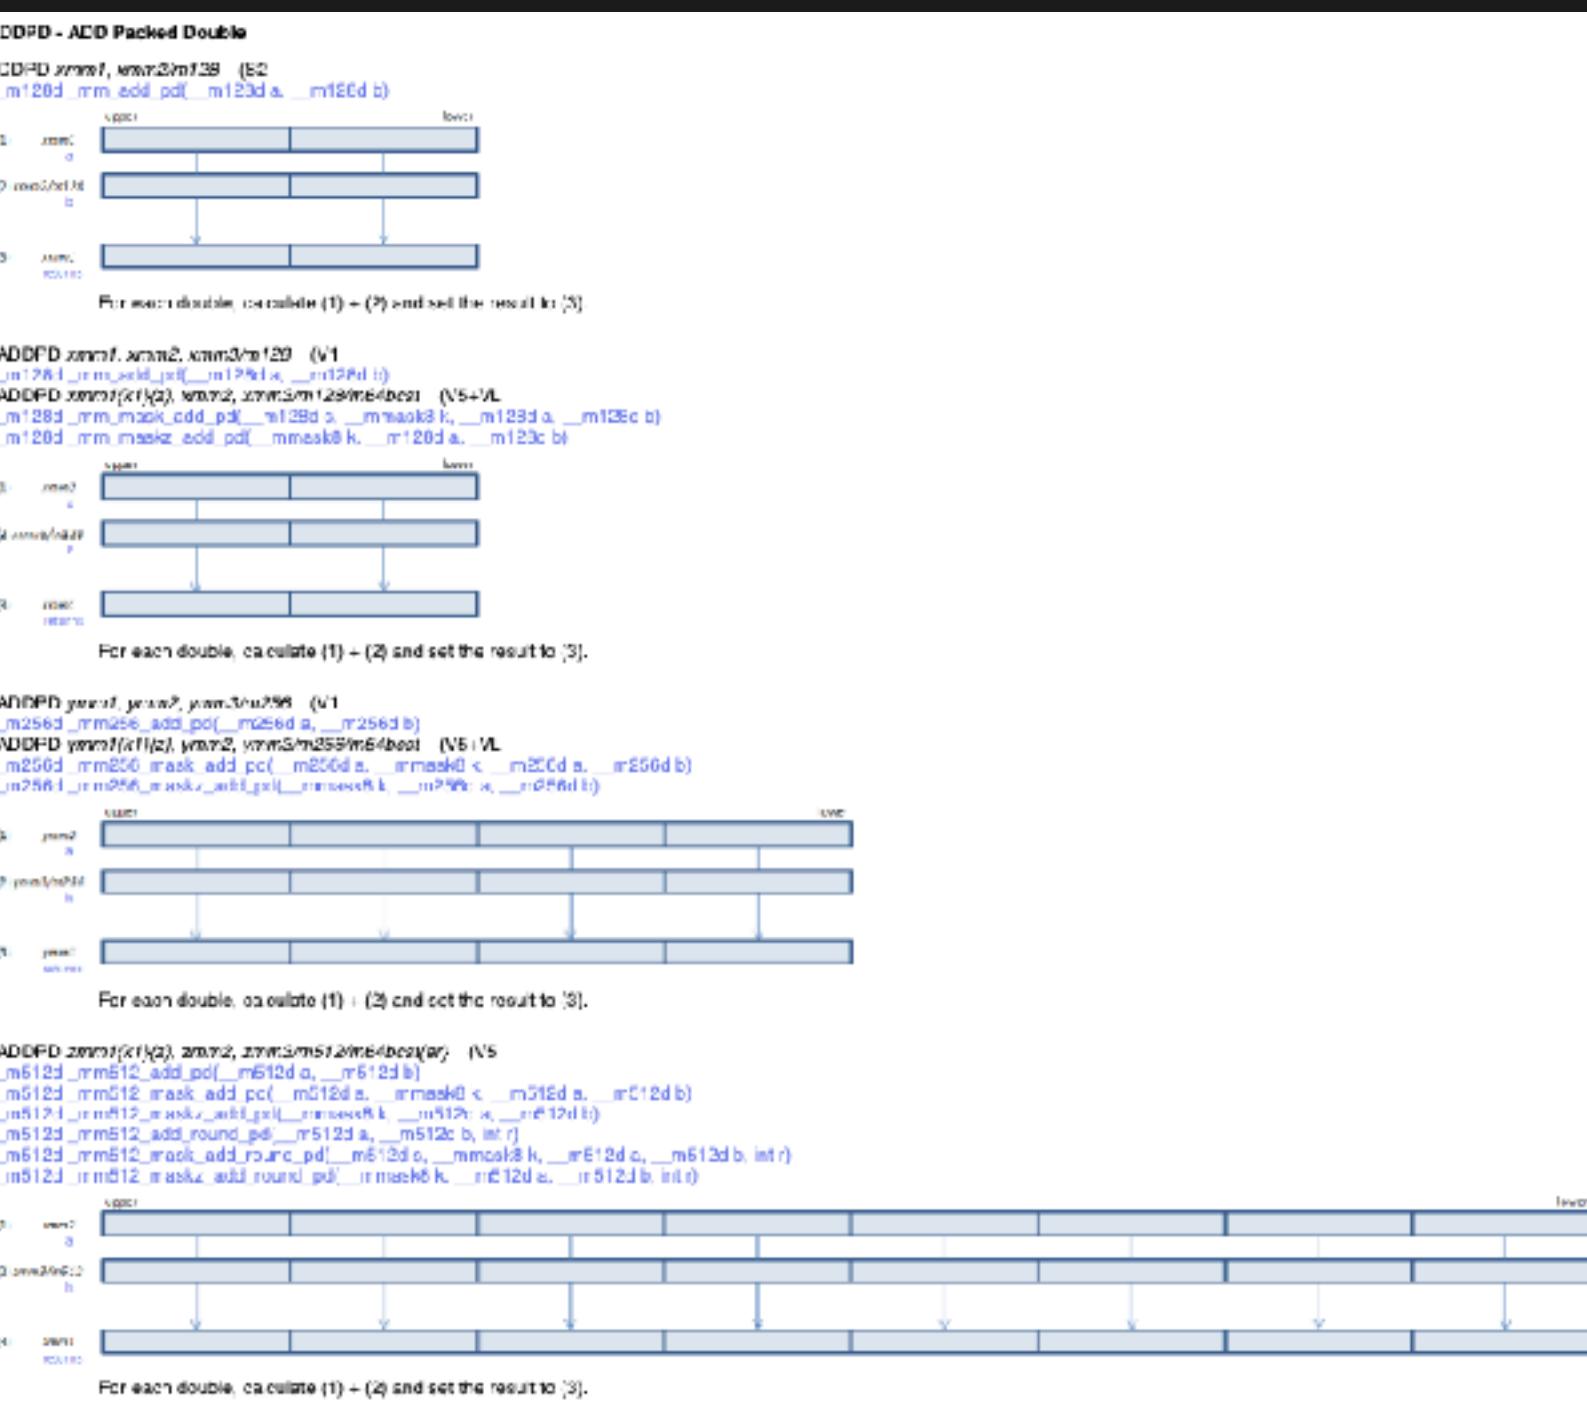


Accessibility: adjust monitor distances

This is the size most subtitles will be

- Most text will be approximately this size
- Some will be larger
- Some slightly smaller



- Code will typically look like this

- Assembly will look like this:

```
fld    QWORD PTR [rsi+rcx*8]
fadd  QWORD PTR [rdx+rcx*8]
fstp  QWORD PTR [rdx+rcx*8]
```

- C and other high level languages will look like this

```
void add(int n, double *x, double *y)
{
    for (int i = 0; i < n; i++)
        y[i] = x[i] + y[i];
}
```

AVX10 for HPC:

A reasonable solution to the 7 levels of AVX-512 folly

Felix LeClair for EasyBuild, Q4 2023

On the docket: Topics covered

1. What is an ISA?
2. x87 “It’s not a great floating point unit, but it’s what we have”
3. MMX “What is SIMD?”
4. SSE 1-2 “SIMD and Real Types”
5. SSE 3-4.2 “”
6. Check-in 1
7. AVX1
8. AVX2
9. Check in 2, electric boogaloo
10. AVX “3” or “I can’t believe it’s not AVX512”
11. AVX “3.5” or “Whatever Phi was”
12. AVX “4” or AVX512 VL, DQ, BW
13. AVX “5” or AVX512 IFMA, VBMI
14. AVX “6” or AVX512BF16
15. AVX “7” or AVX512 VPOPCNTDQ, VNNI, VBMI2, BITALG
16. AVX “8” or VP2INTERSECT
17. AVX “9” or FP16
18. AVX “10” or “The One we care about”

**The ISA is the language your computer speaks.
Just like natural languages, it can have several
dialects (variants) and regional terms
(extensions) only some computers understand.**

Robert Clausecker

**For programmers:
ISA : CPU :: Language
standard : compiler**

Alexander Monakov

**The ISA is the
language the chip
understands.**

CHAOSHACKER

**[An ISA] is a mapping
between integers
and operations**

Violet the Chromatic Bat

**An ISA is like musical notes, [it
makes] an orchestra of
transistors perform a
symphony!**

Andreas Herten

**[An ISA] is like a box
of chocolates**

Brandon Biggs

**[An ISA] is like a
vocabulary[.]**

Moritz Lehmann

[An ISA] is a contract between hardware and software designers that lets us work together

Andrew Richards

For higher-level programmers, [an ISA] is the stable API for the CPU.

Andrew Richards

**An ISA is the train-line to which HW designers
chain their screaming future colleagues to.**

Tom Forsyth

Section 1

Floating point OR SIMD

An ISA from 1978 isn't going to fit the needs of today

"I heard you like extensions, so I put extensions extending your extended extensions"

- Much like Fortran 77 Code today, only implementing HW specs from the 70's can be loosely described as a poor life choice
 - The original 8086 and 8088 didn't have HW support for Floating point.
 - Enjoy your $\sim O(N^2)$ performance whenever you need to divide A/B

The 8086

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);

}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

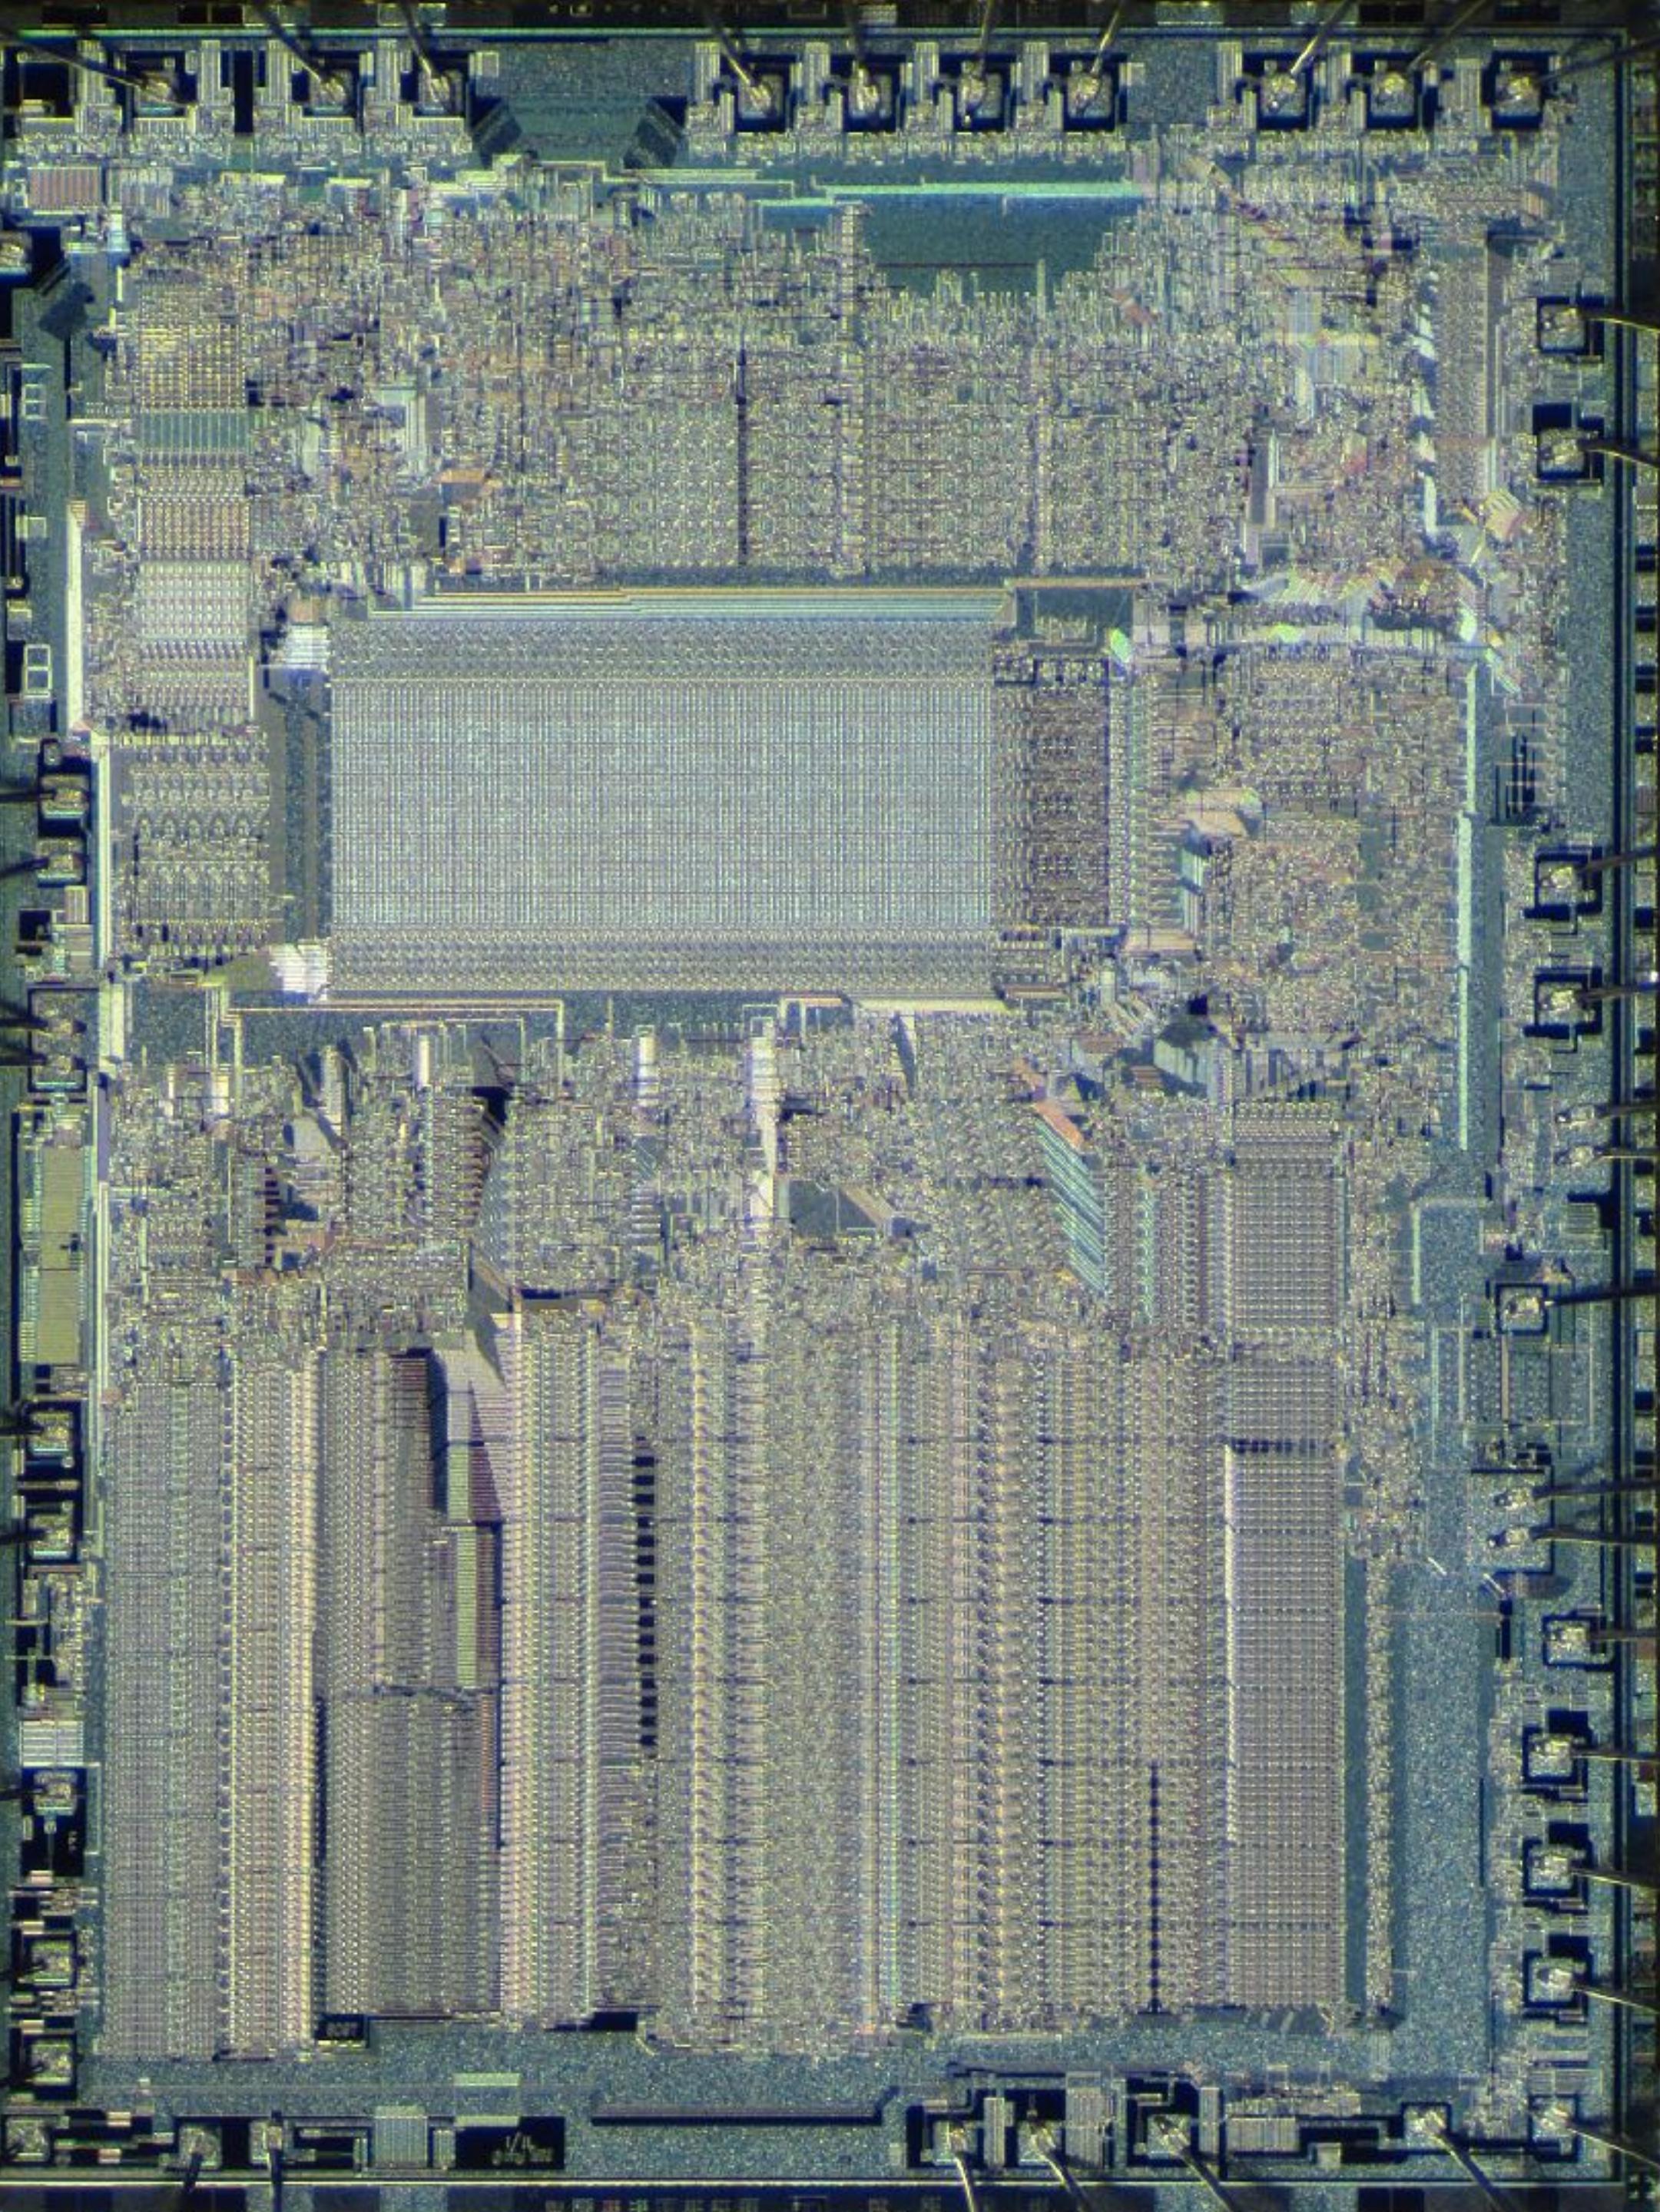
    return 0;
}
```

```
DAXPY(int, double*, double*, double*):
    push    ebp
    push    ebx
    push    edi
    push    esi
    sub     esp,0xc
    call    c <DAXPY(int, double*, double*, double*)+0xc>
    pop     ebx
    add    ebx,0x3
    R_386_GOTPC
_GLOBAL_OFFSET_TABLE_
    sub     esp,0xc
    mov     ebp,0x2000
    push    ebp
    call    7c <main+0x1d>
    R_386_PLT32 operator new[]
(unsigned int)
    add    esp,0x10
    mov     esi,eax
    sub     esp,0xc
    push    ebp
    call    8a <main+0x2b>
    R_386_PLT32 operator new[]
(unsigned int)
    add    esp,0x10
    mov     edi,eax
    sub     esp,0xc
    push    ebp
    call    98 <main+0x39>
    R_386_PLT32 operator new[]
(unsigned int)
    add    esp,0x10
    push   eax
    push   edi
    push   esi
    push   0x400
    call    a8 <main+0x49>
    R_386_PC32 DAXPY(int, double*, double*, double*)
    add    esp,0x10
    xor    eax,eax
    add    esp,0xc
    pop     esi
    pop     edi
    pop     ebx
    pop     ebp
    ret
main:
```

The 8087-1980

Floating point on x86!

- 3 data types:
 - FP32
 - FP64
 - FP80
- Inspired what would become the IEEE-754 (1985) floating spec we know and love
- 80bit registers



The 8087

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);

}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

    return 0;
}
```

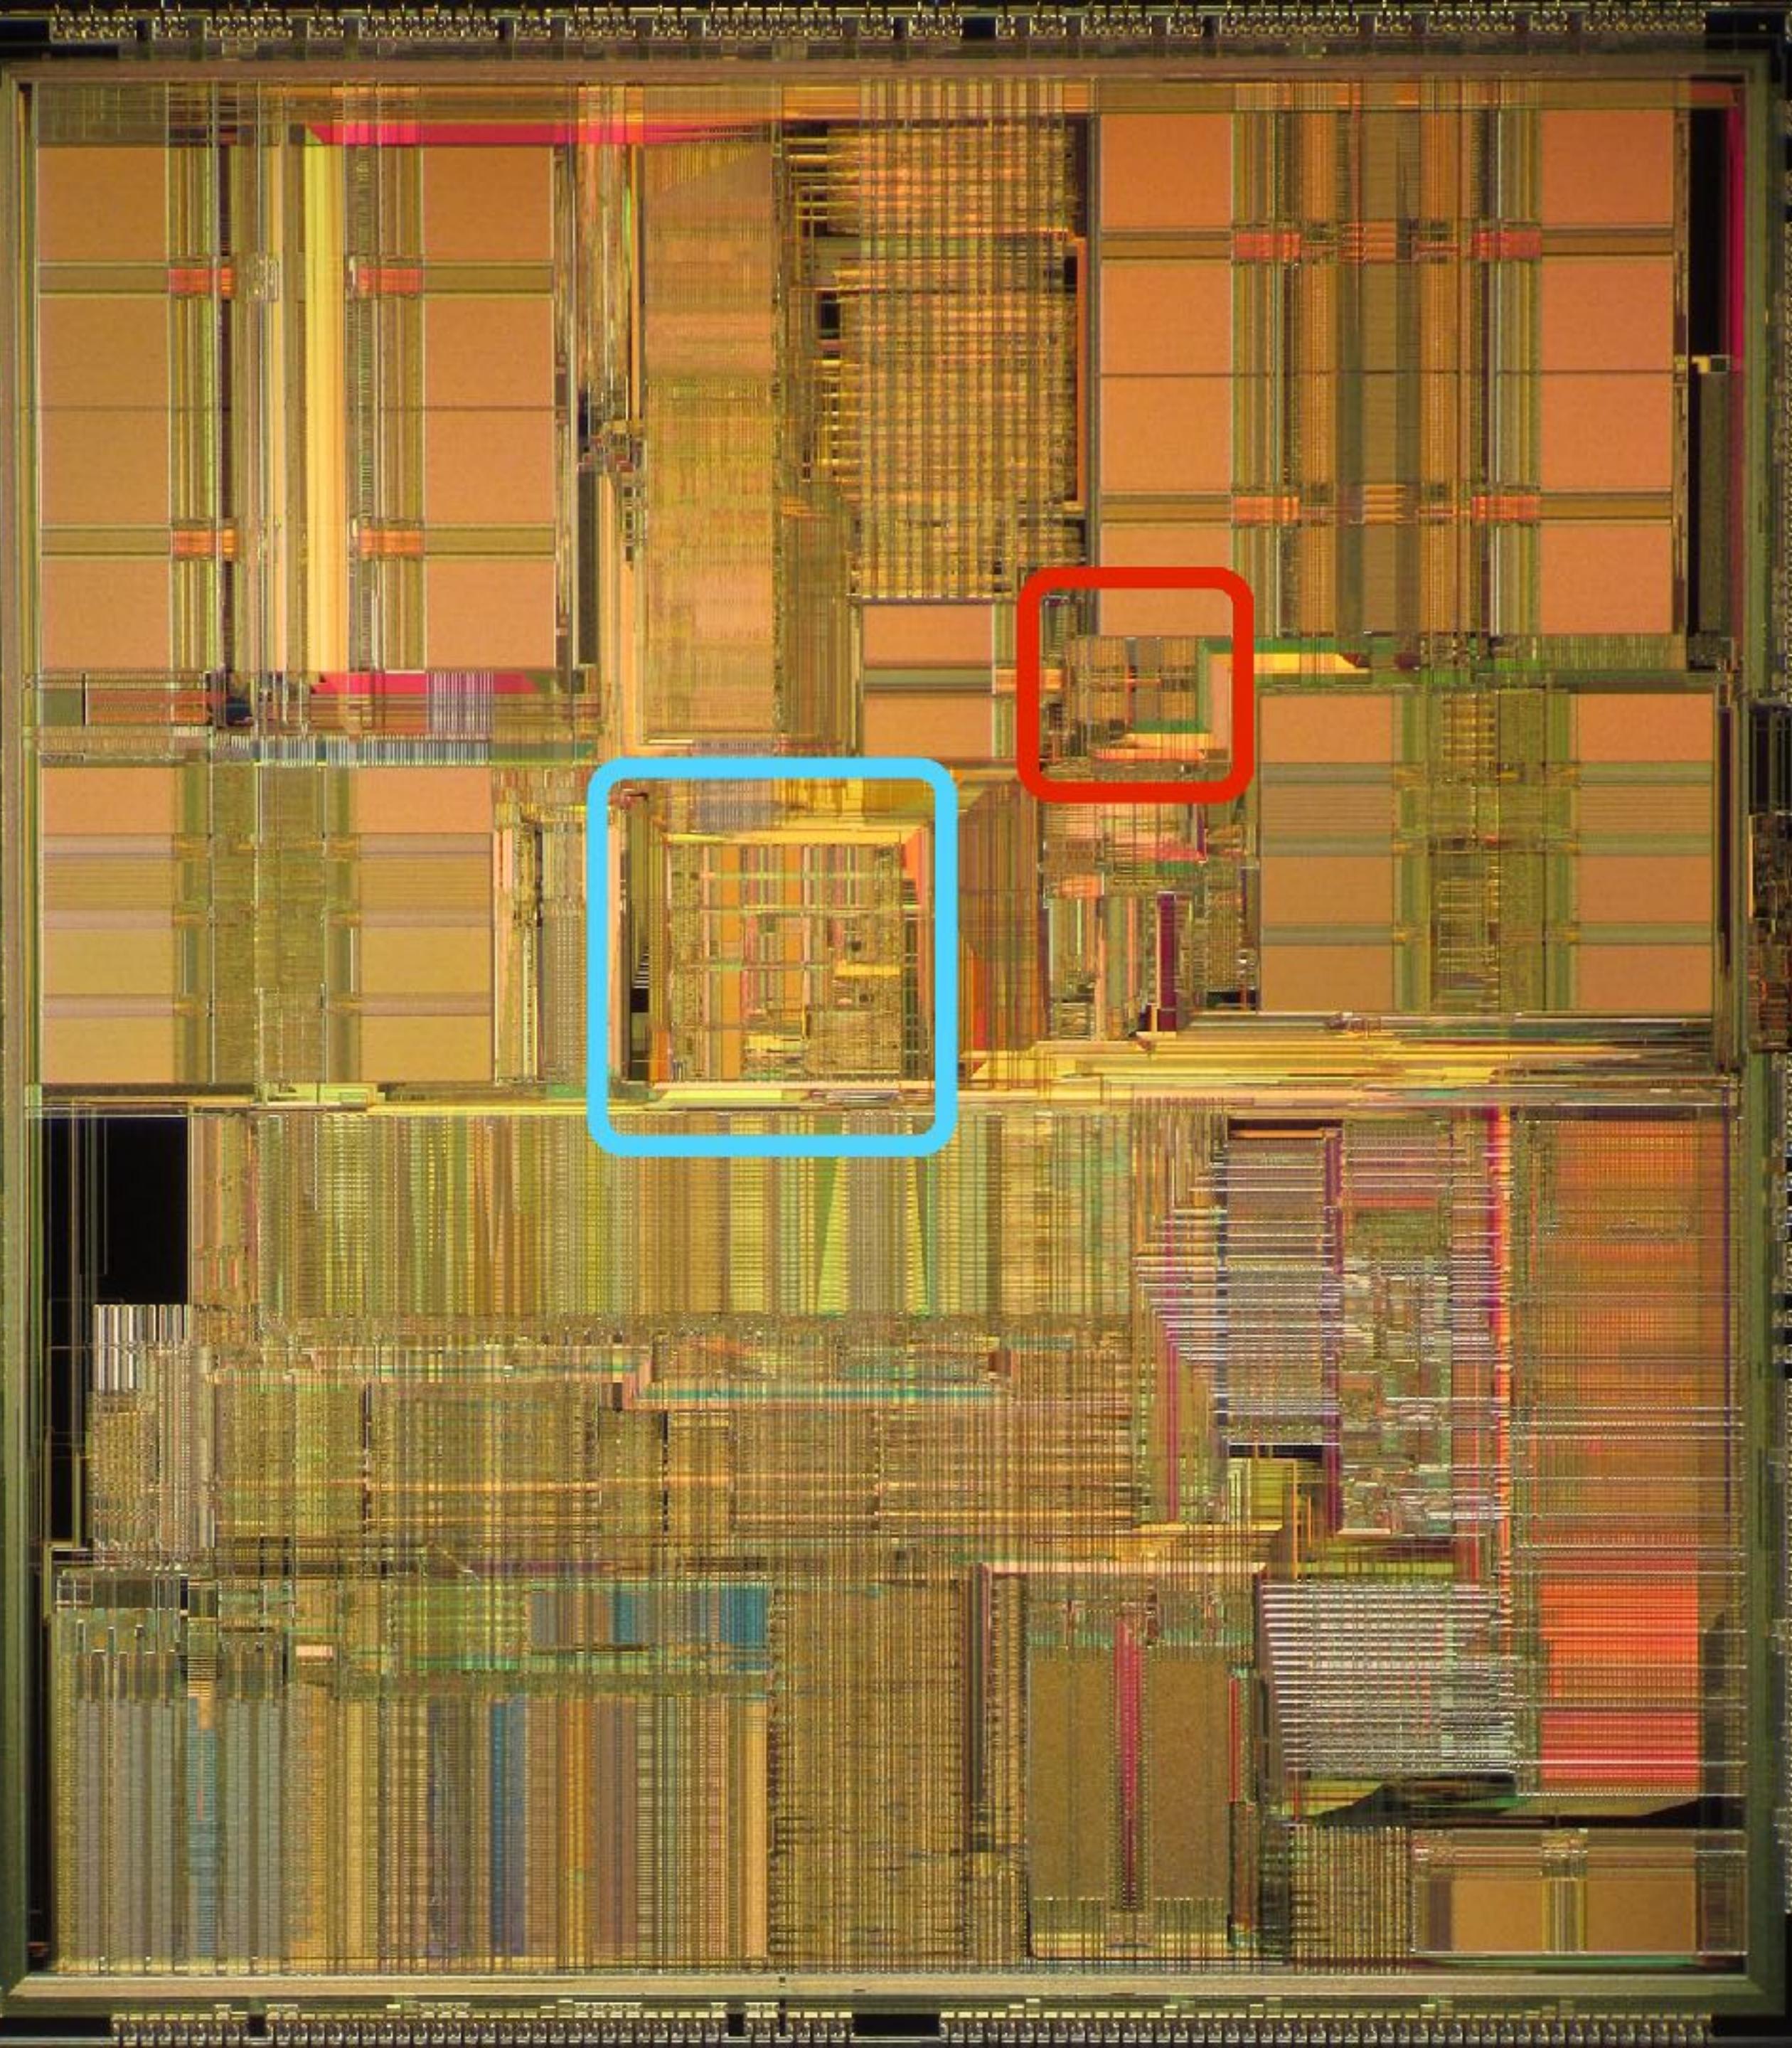
```
DAXPY(int, double*, double*, double*):
    push    edi
    push    esi
    xor     eax,eax
    mov     ecx,WORD PTR [esp+0x18]
    mov     edx,WORD PTR [esp+0x14]
    mov     esi,WORD PTR [esp+0x10]
    mov     edi,WORD PTR [esp+0xc]
    fld     QWORD PTR [esi+eax*8]
    fadd   QWORD PTR [edx+eax*8]
    fmul   QWORD PTR [ecx+eax*8]
    fstp   QWORD PTR [ecx+eax*8]
    inc    eax
    cmp    eax,edi
    jl    14 <DAXPY(int,
double*, double*, double*)+0x14>
    pop    esi
    pop    edi
    ret
main:
    push    ebp
    push    ebx
    push    edi
    push    esi
    sub    esp,0xc
    call    34 <main+0xc>
    pop    ebx
    add    ebx,0x3
    R_386_GOTPC
_GLOBAL_OFFSET_TABLE_
```

```
sub    esp,0xc
mov    ebp,0x2000
push   ebp
call   45 <main+0x1d>
    R_386_PLT32 operator
new[ ](unsigned int)
    add    esp,0x10
    mov    esi,eax
    sub    esp,0xc
    push   ebp
    call   53 <main+0x2b>
    R_386_PLT32 operator
new[ ](unsigned int)
    add    esp,0x10
    mov    edi,eax
    sub    esp,0xc
    push   ebp
    call   61 <main+0x39>
    R_386_PLT32 operator
new[ ](unsigned int)
    add    esp,0x10
    push  eax
    push  edi
    push  esi
    push  0x400
    call  71 <main+0x49>
    R_386_PC32 DAXPY(int,
double*, double*, double*)
    add    esp,0x10
    xor    eax,eax
    add    esp,0xc
    pop    esi
    pop    edi
    pop    ebx
    pop    ebp
    ret
```

MMX -1997

SIMD on x86

- 4 data types:
 - 1* Int64
 - 2* Int32
 - 4* Int16
 - 8* Int8
- Used the lower 64 bits of the 80 bit x87 register



MMX

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);

}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

    return 0;
}
```

```
DAXPY(int, double*, double*, double*):
    push    edi
    push    esi
    xor     eax,eax
    mov     ecx,WORD PTR [esp+0x18]
    mov     edx,WORD PTR [esp+0x14]
    mov     esi,WORD PTR [esp+0x10]
    mov     edi,WORD PTR [esp+0xc]
    fld     QWORD PTR [esi+eax*8]
    fadd   QWORD PTR [edx+eax*8]
    fmul   QWORD PTR [ecx+eax*8]
    fstp   QWORD PTR [ecx+eax*8]
    inc    eax
    cmp    eax,edi
    jl    14 <DAXPY(int,
double*, double*, double*)+0x14>
    pop    esi
    pop    edi
    ret
main:
    push    ebp
    push    ebx
    push    edi
    push    esi
    sub    esp,0xc
    call    34 <main+0xc>
    pop    ebx
    add    ebx,0x3
    R_386_GOTPC
_GLOBAL_OFFSET_TABLE_
```

sub esp,0xc mov ebp,0x2000 push ebp call 45 <main+0x1d> R_386_PLT32 operator new[](unsigned int) add esp,0x10 mov esi,eax sub esp,0xc push ebp call 53 <main+0x2b> R_386_PLT32 operator new[](unsigned int) add esp,0x10 mov edi,eax sub esp,0xc push ebp call 61 <main+0x39> R_386_PLT32 operator new[](unsigned int) add esp,0x10 push eax push edi push esi push 0x400 call 71 <main+0x49> R_386_PC32 DAXPY(int, double*, double*, double*) add esp,0x10 xor eax,eax add esp,0xc pop esi pop edi pop ebx pop ebp ret
--

X87

Perform DAXPY z[]=(x[]+y[])*z[]

```

DAXPY(int, double*, double*,
double*):
    push  edi
    push  esi
    xor   eax,eax
    mov   ecx,WORD PTR [esp+0x18]
    mov   edx,WORD PTR [esp+0x14]
    mov   esi,WORD PTR [esp+0x10]
    mov   edi,WORD PTR [esp+0xc]
    fld   QWORD PTR [esi+eax*8]
    fadd  QWORD PTR [edx+eax*8]
    fmul  QWORD PTR [ecx+eax*8]
    fstp  QWORD PTR [ecx+eax*8]
    inc   eax
    cmp   eax,edi
    jl    14 <DAXPY(int, double*,
double*, double*)+0x14>
    pop   esi
    pop   edi
    ret
main:
    push  ebp
    push  ebx
    push  edi
    push  esi
    sub   esp,0xc
    call  34 <main+0xc>
    pop   ebx
    add   ebx,0x3
    R_386_GOTPC
_GLOBAL_OFFSET_TABLE_
    sub   esp,0xc
    mov   ebp,0x2000
    push  ebp
    call  45 <main+0x1d>

```

```

        R_386_PLT32 operator new[ ]
(unsigned int)
    add   esp,0x10
    mov   esi,eax
    sub   esp,0xc
    push  ebp
    call  53 <main+0x2b>
        R_386_PLT32 operator new[ ]
(unsigned int)
    add   esp,0x10
    mov   edi,eax
    sub   esp,0xc
    push  ebp
    call  61 <main+0x39>
        R_386_PLT32 operator new[ ]
(unsigned int)
    add   esp,0x10
    push  eax
    push  edi
    push  esi
    push  0x400
    call  71 <main+0x49>
        R_386_PC32 DAXPY(int,
double*, double*, double*)
    add   esp,0x10
    xor   eax,eax
    add   esp,0xc
    pop   esi
    pop   edi
    pop   ebx
    add   ebx,0x3
    R_386_GOTPC
_GLOBAL_OFFSET_TABLE_
    sub   esp,0xc
    mov   ebp,0x2000
    push  ebp
    ret

```

```

DAXPY(int, double*, double*,
double*):
    push  edi
    push  esi
    xor   eax,eax
    mov   ecx,WORD PTR [esp+0x18]
    mov   edx,WORD PTR [esp+0x14]
    mov   esi,WORD PTR [esp+0x10]
    mov   edi,WORD PTR [esp+0xc]
    fld   QWORD PTR [esi+eax*8]
    fadd  QWORD PTR [edx+eax*8]
    fmul  QWORD PTR [ecx+eax*8]
    fstp  QWORD PTR [ecx+eax*8]
    inc   eax
    cmp   eax,edi
    jl    14 <DAXPY(int, double*,
double*, double*)+0x14>
    pop   esi
    pop   edi
    ret
main:
    push  ebp
    push  ebx
    push  edi
    push  esi
    sub   esp,0xc
    call  34 <main+0xc>
    pop   ebx
    add   ebx,0x3
    R_386_GOTPC
_GLOBAL_OFFSET_TABLE_
    sub   esp,0xc
    mov   ebp,0x2000
    push  ebp
    call  45 <main+0x1d>

```

```

        R_386_PLT32 operator new[ ]
(unsigned int)
    add   esp,0x10
    mov   esi,eax
    sub   esp,0xc
    push  ebp
    call  53 <main+0x2b>
        R_386_PLT32 operator new[ ]
(unsigned int)
    add   esp,0x10
    mov   edi,eax
    sub   esp,0xc
    push  ebp
    call  61 <main+0x39>
        R_386_PLT32 operator new[ ]
(unsigned int)
    add   esp,0x10
    push  eax
    push  edi
    push  esi
    push  0x400
    call  71 <main+0x49>
        R_386_PC32 DAXPY(int,
double*, double*, double*)
    add   esp,0x10
    xor   eax,eax
    add   esp,0xc
    pop   esi
    pop   edi
    pop   ebx
    pop   ebp
    ret

```

8087

Perform DAXPY $z[] = (x[] + y[]) * z[]$

MMX

Perform DAXPY $z[] = (x[] + y[]) * z[]$



Section 2

Floating point AND SIMD

SSE1-1999

Streaming SIMD extensions

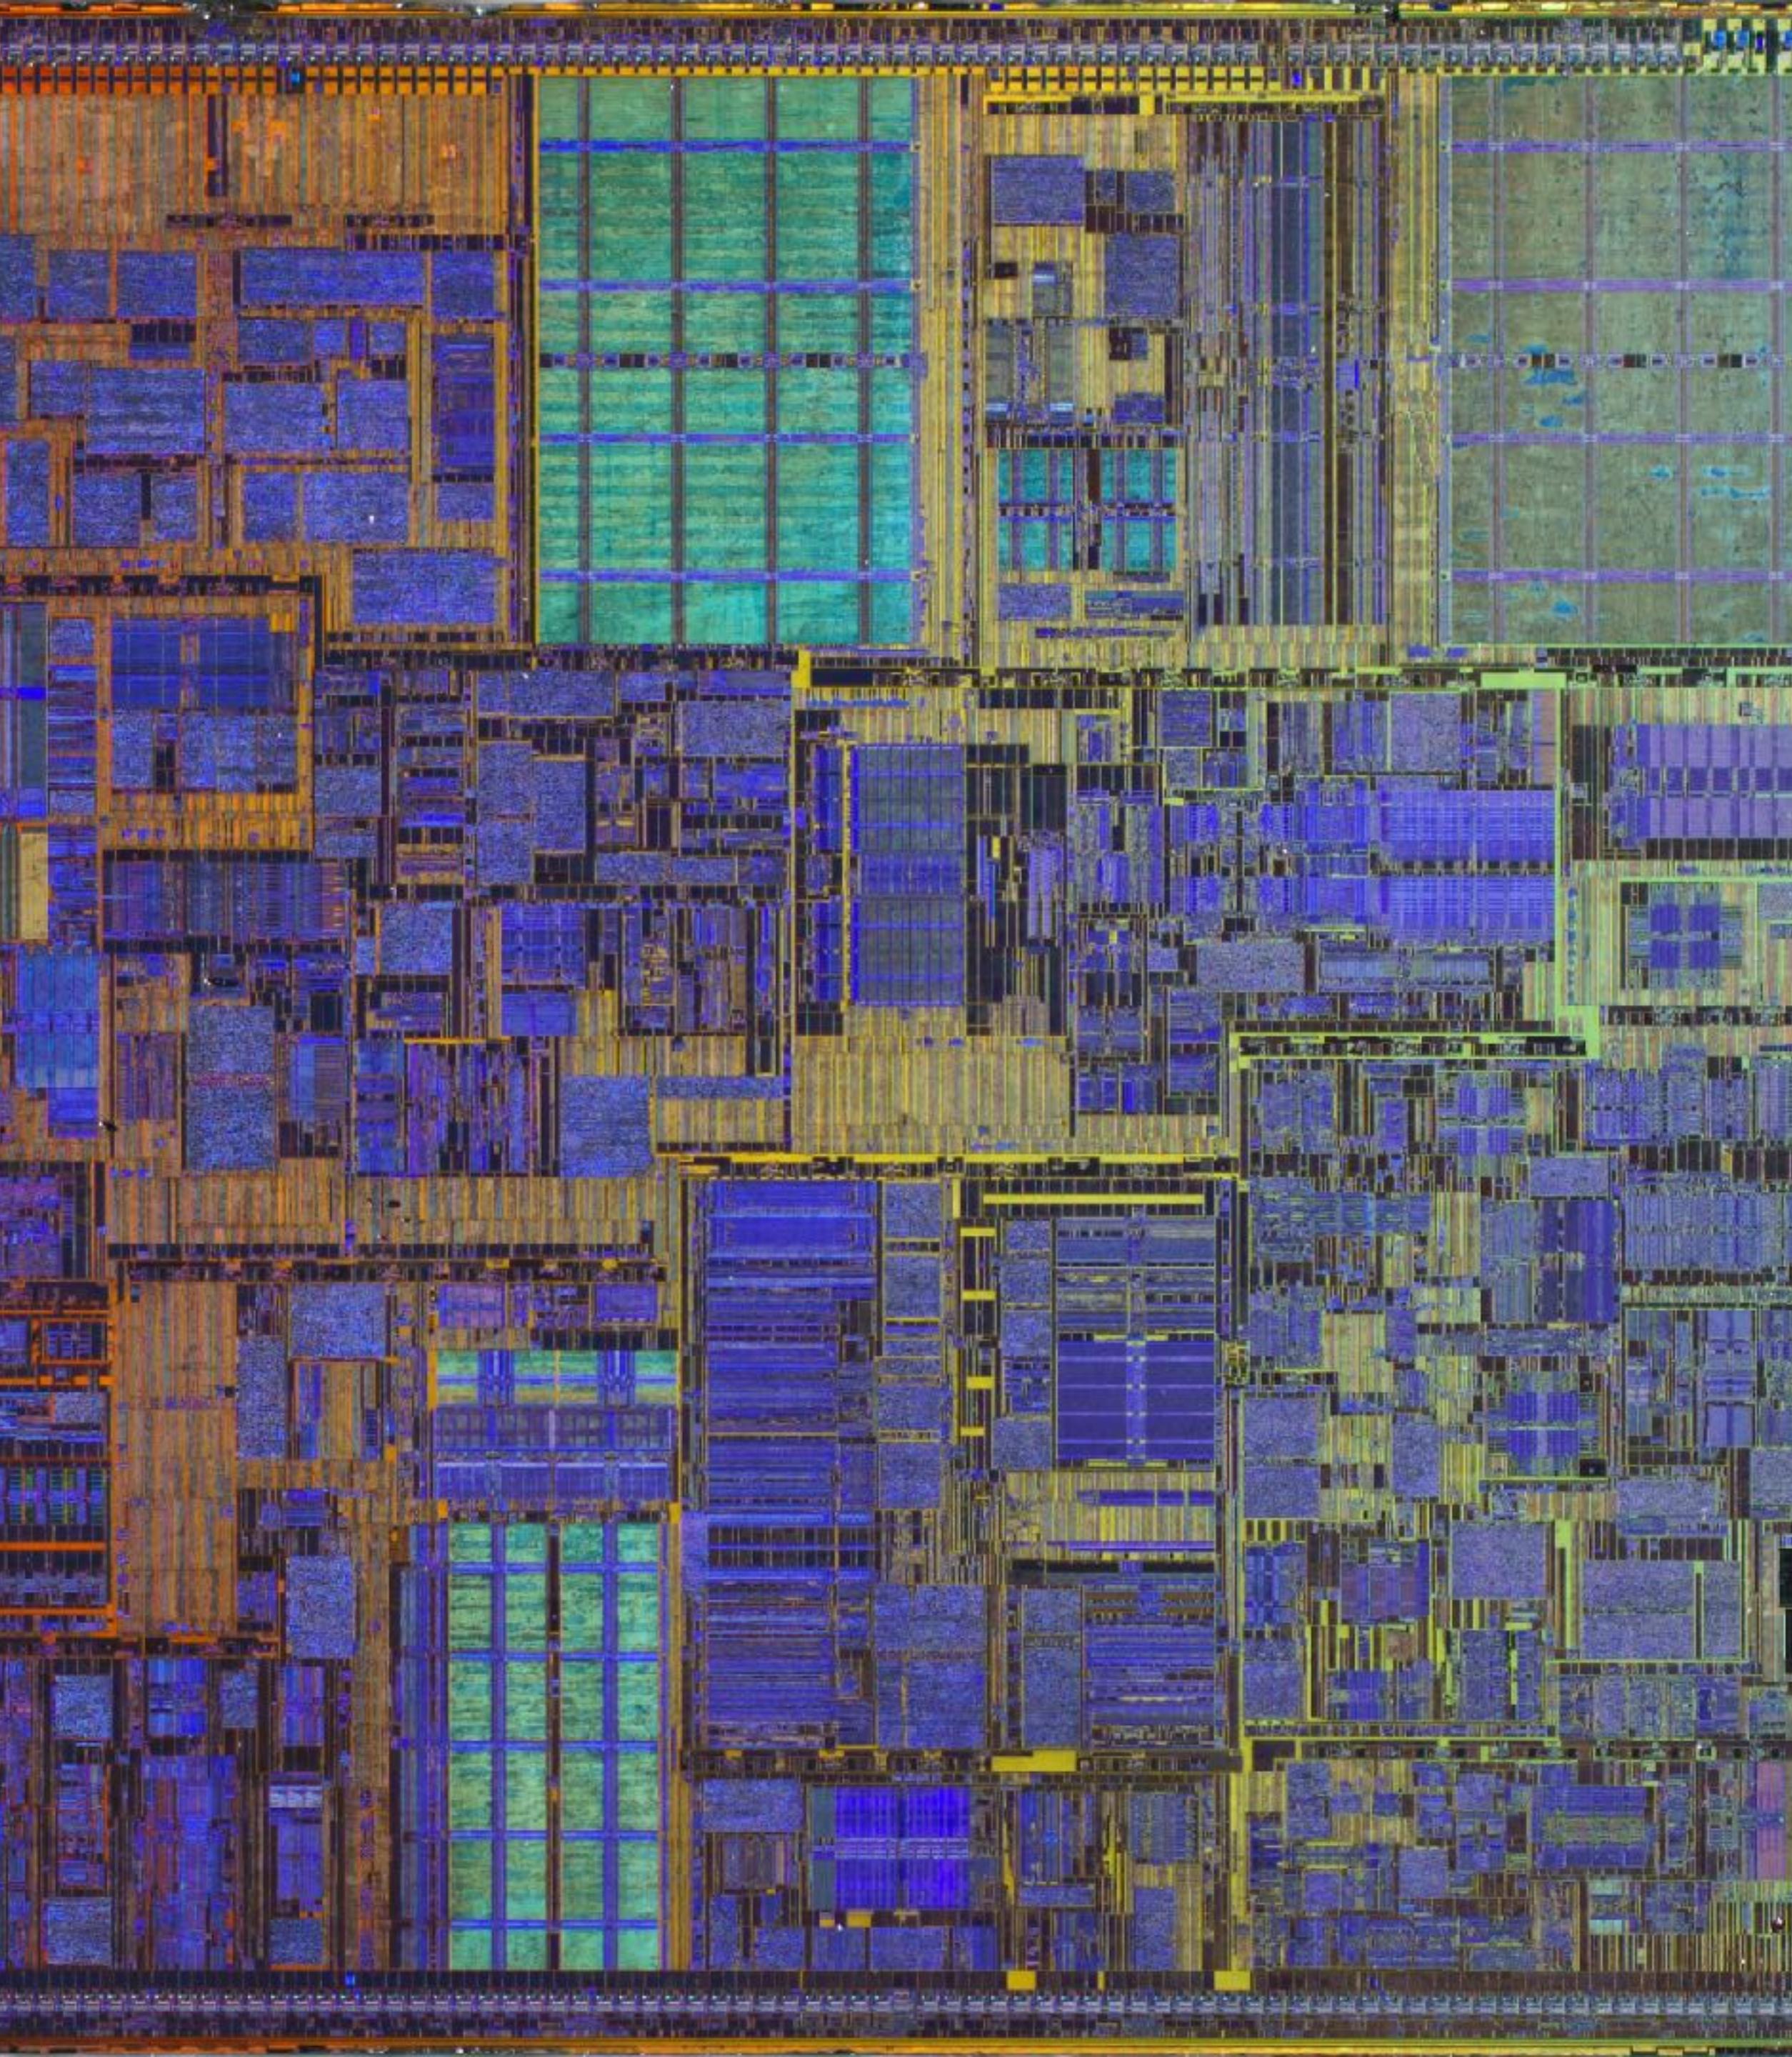
- 1 data type:
 - 4* FP32
- 8 new registers
 - xmm0-7
 - 128 bits wide



SSE2-2000

“Real” general purpose SIMD

- 6 data types:
 - 2* FP64(!!!)
 - 4* FP32
 - 2* Int64
 - 4* Int32
 - 8* Int16
 - 16* Int8



SSE2

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);
}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

    return 0;
}
```

```
DAXPY(int, double*, double*, double*):
    cmp    edi,0x2
    mov    eax,0x1
    cmovge eax,edi
    cmp    eax,0xa
    jae    17 <DAXPY(int,
double*, double*, double*)
+0x17>
    xor    edi,edi
    jmp    b6 <DAXPY(int,
double*, double*, double*)
+0xb6>
    lea    rdi,[rcx+rax*8]
    lea    r8,[rsi+rax*8]
    lea    r9,[rdx+rax*8]
    cmp    r8,rcx
    seta   r10b
    cmp    rdi,rsi
    seta   r11b
    cmp    r9,rcx
    seta   r8b
    cmp    rdi,rdx
    seta   r9b
    xor    edi,edi
    test   r10b,r11b
    jne    b6 <DAXPY(int,
double*, double*, double*)
+0xb6>
    and    r8b,r9b
    jne    b6 <DAXPY(int,
double*, double*, double*)
+0xb6>
    mov    edi,eax
    and    edi,0x7fffffff
    xor    r8d,r8d
    cs    nop WORD PTR
[rax+rax*1+0x0]
    movupd xmm0,XMMWORD PTR
[rsi+r8*8]
    movupd xmm1,XMMWORD PTR
[rsi+r8*8+0x10]
    movupd xmm2,XMMWORD PTR
[rdx+r8*8]
    addpd xmm2,xmm0
    movupd xmm0,XMMWORD PTR
[rdx+r8*8+0x10]
    addpd xmm0,xmm1
    movupd xmm1,XMMWORD PTR
[rcx+r8*8]
    mulpd xmm1,xmm2
    movupd xmm2,XMMWORD PTR
[rcx+r8*8+0x10]
    mulpd xmm2,xmm0
    movupd XMMWORD PTR
[rcx+r8*8],xmm1
    movupd XMMWORD PTR
[rcx+r8*8+0x10],xmm2
    add    r8,0x4
    cmp    rdi,r8
    jne    60 <DAXPY(int,
double*, double*, double*)
+0x60>
    cmp    rdi,rax
    je     165 <DAXPY(int,
double*, double*, double*)
+0x165>
    mov    r8,rdi
    not    r8
    add    r8,rax
    mov    r9,rax
    and    r9,0x3
    je     ec <DAXPY(int,
double*, double*, double*)
+0xec>
    nop    DWORD PTR
[rax+rax*1+0x0]
    movsd xmm0,QWORD PTR
[rsi+rdi*8]
    addsd xmm0,QWORD PTR
[rdx+rdi*8]
    mulsd xmm0,QWORD PTR
[rcx+rdi*8]
    movsd QWORD PTR
[rcx+rdi*8],xmm0
    inc    rdi
    dec    r9
    jne    d0 <DAXPY(int,
double*, double*, double*)
+0xd0>
    cmp    r8,0x3
    jb    165 <DAXPY(int,
double*, double*, double*)
+0x165>
    data16 data16 data16 data16
    cs    nop WORD PTR
[rax+rax*1+0x0]
    movsd xmm0,QWORD PTR
[rsi+r8*8]
    addsd xmm0,QWORD PTR
[rdx+rdi*8]
    mulsd xmm0,QWORD PTR
[rcx+rdi*8]
    movsd QWORD PTR
[rcx+rdi*8],xmm0
    movsd xmm0,QWORD PTR
[rsi+rdi*8+0x8]
    addsd xmm0,QWORD PTR
[rdx+rdi*8+0x8]
    mulsd xmm0,QWORD PTR
[rcx+rdi*8+0x8]
    movsd QWORD PTR
[rcx+rdi*8+0x8],xmm0
    movsd xmm0,QWORD PTR
[rsi+rdi*8+0x10]
    addsd xmm0,QWORD PTR
[rdx+rdi*8+0x10]
    add    r8,0x4
    mulsd xmm0,QWORD PTR
[rcx+rdi*8+0x10]
    movsd xmm0,QWORD PTR
[rsi+rdi*8+0x18]
    addsd xmm0,QWORD PTR
[rdx+rdi*8+0x18]
    mulsd xmm0,QWORD PTR
[rcx+rdi*8+0x18]
    movsd QWORD PTR
[rcx+rdi*8+0x18],xmm0
    add    rdi,0x4
    cmp    rax,rdi
    jne    100 <DAXPY(int,
double*, double*, double*)
+0x100>
    ret
    cs    nop WORD PTR
[rax+rax*1+0x0]
main:
    push   r14
    push   rbx
    push   rax
    mov    edi,0x2000
    call   17e <main+0xe>
    R_X86_64_PLT32 operator
new[](unsigned long)-0x4
    mov    rbx,rax
    mov    edi,0x2000
    call   18b <main+0x1b>
    R_X86_64_PLT32 operator
new[](unsigned long)-0x4
    mov    r14,rax
    mov    edi,0x2000
    call   198 <main+0x28>
    R_X86_64_PLT32 operator
new[](unsigned long)-0x4
    mov    edi,0x400
    mov    rsi,rbx
    mov    rdx,r14
    mov    rcx,rax
    call   1ab <main+0x3b>
    R_X86_64_PLT32 DAXPY(int,
double*, double*, double*)-0x4
    xor    eax,eax
    add    rsp,0x8
    pop    rbx
    pop    r14
    ret
```

SSE2 IS

STILL

The default ISA target

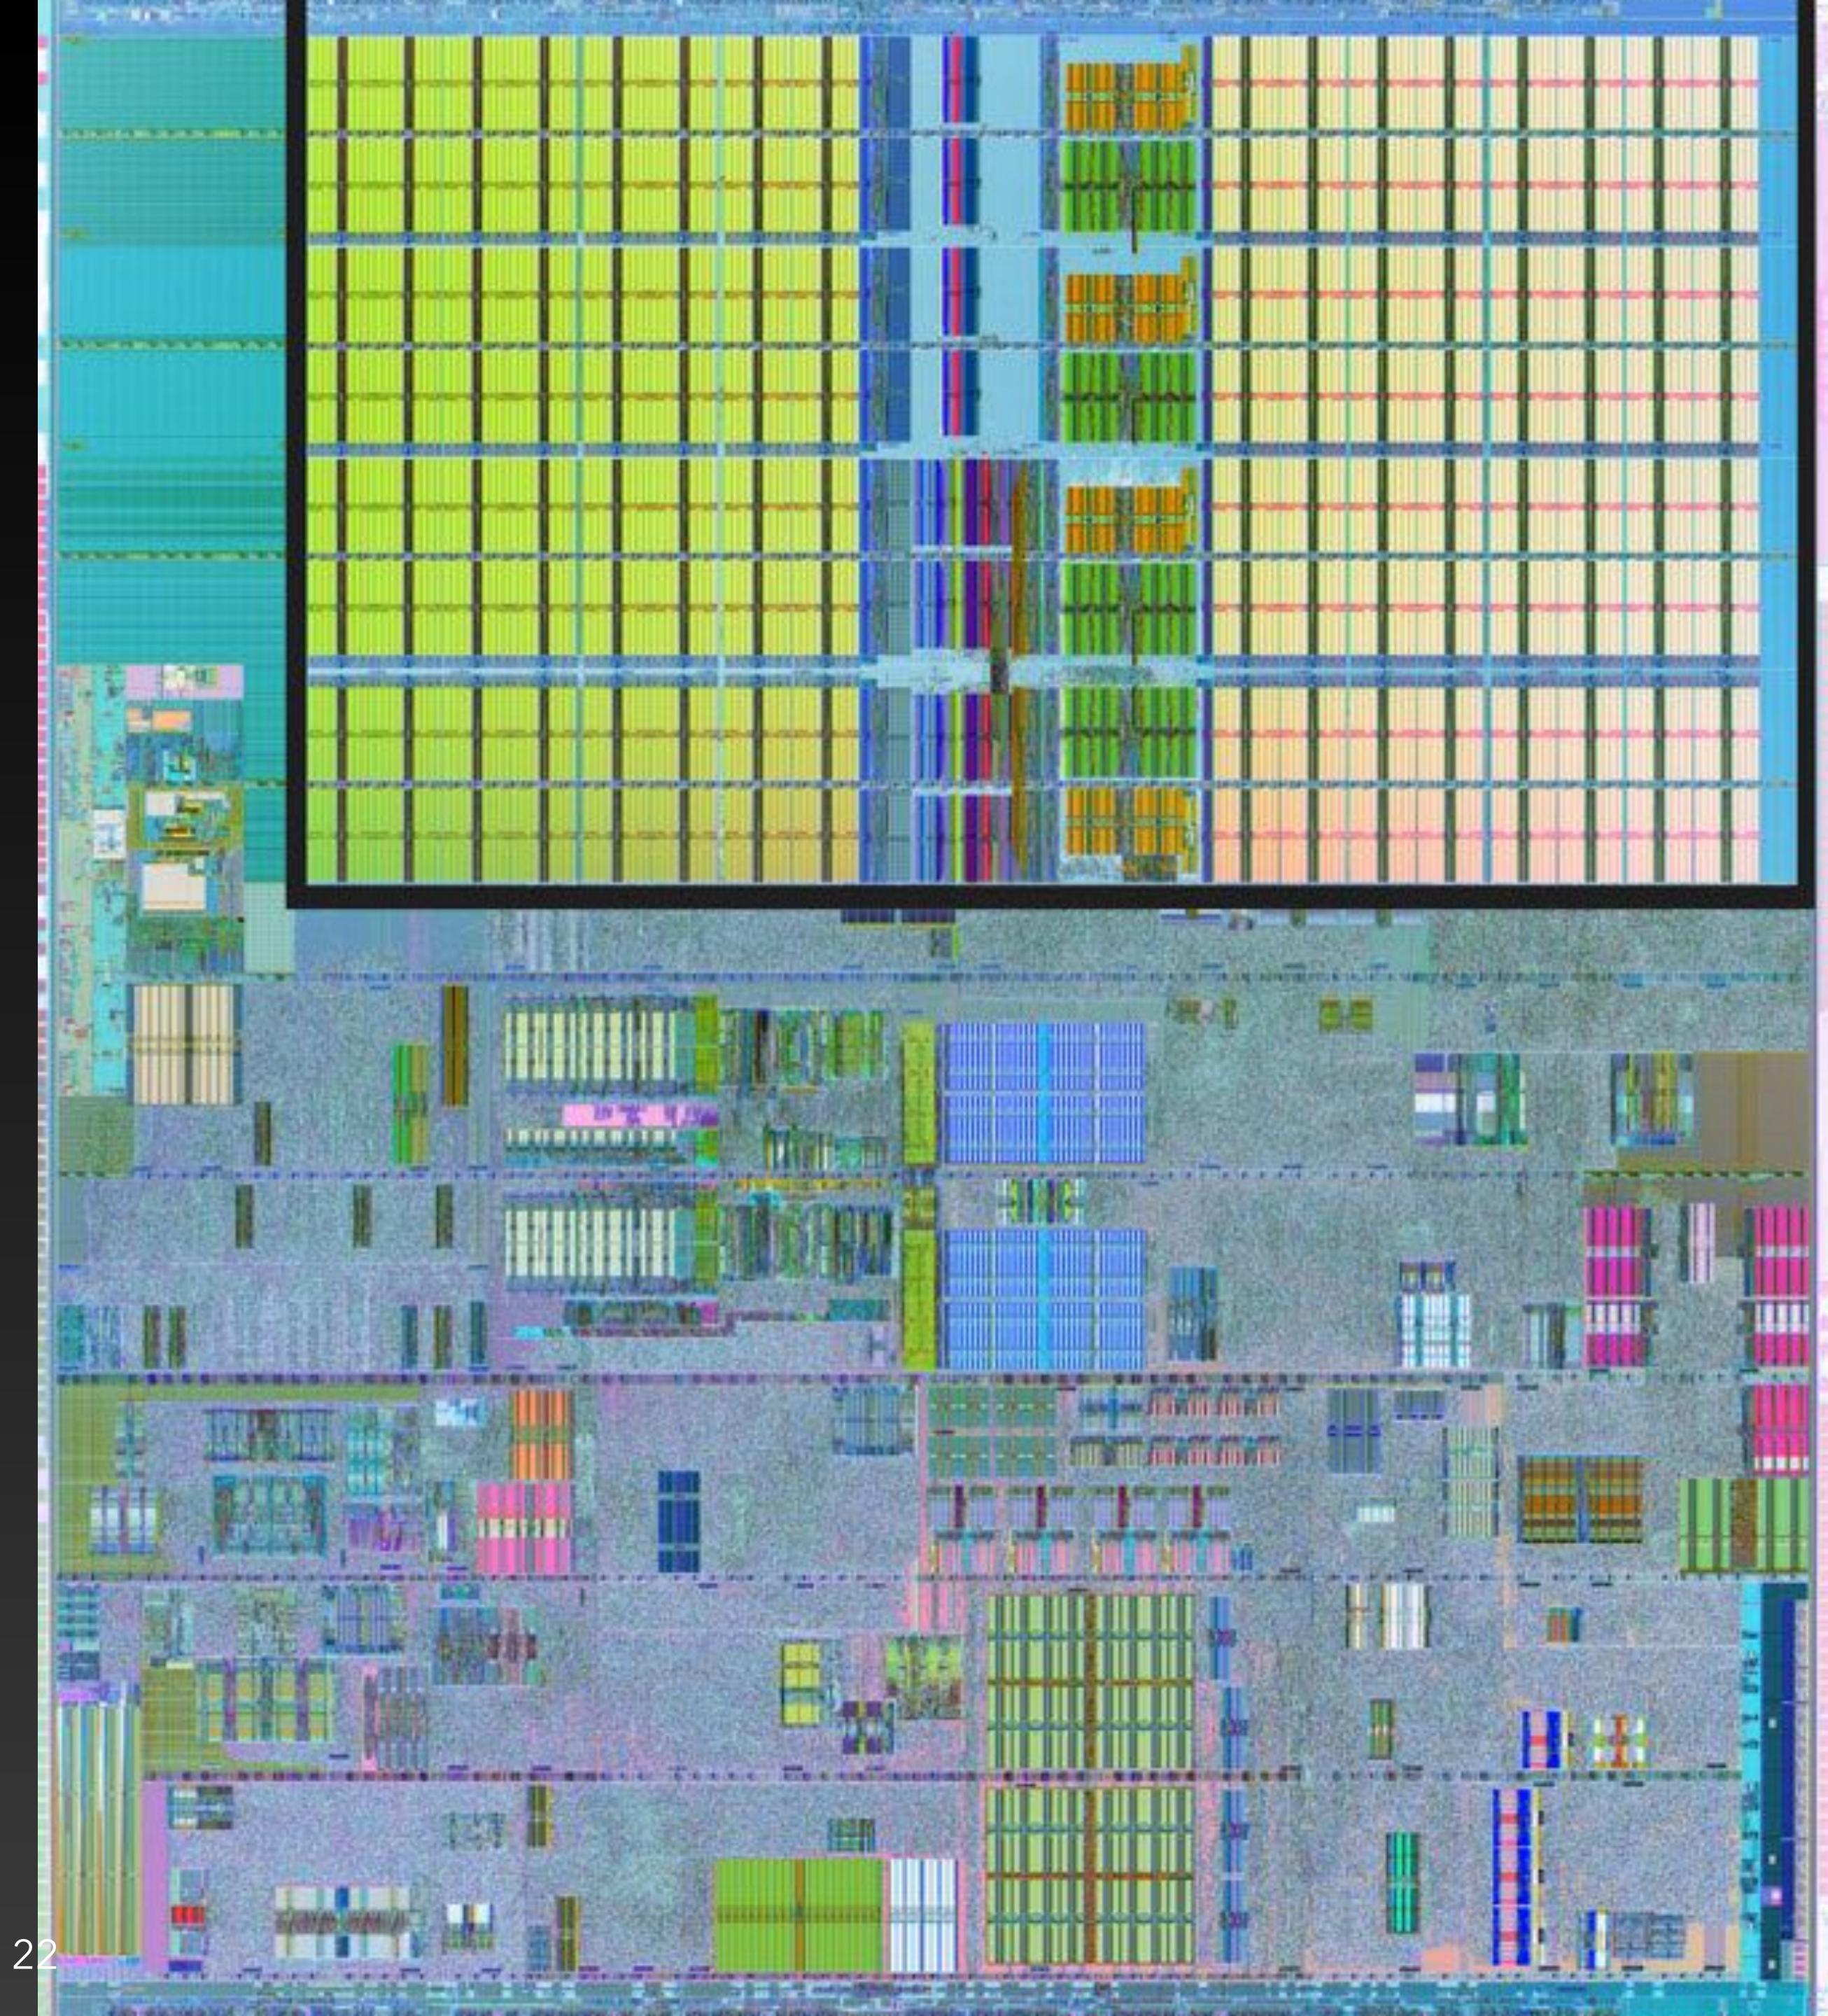
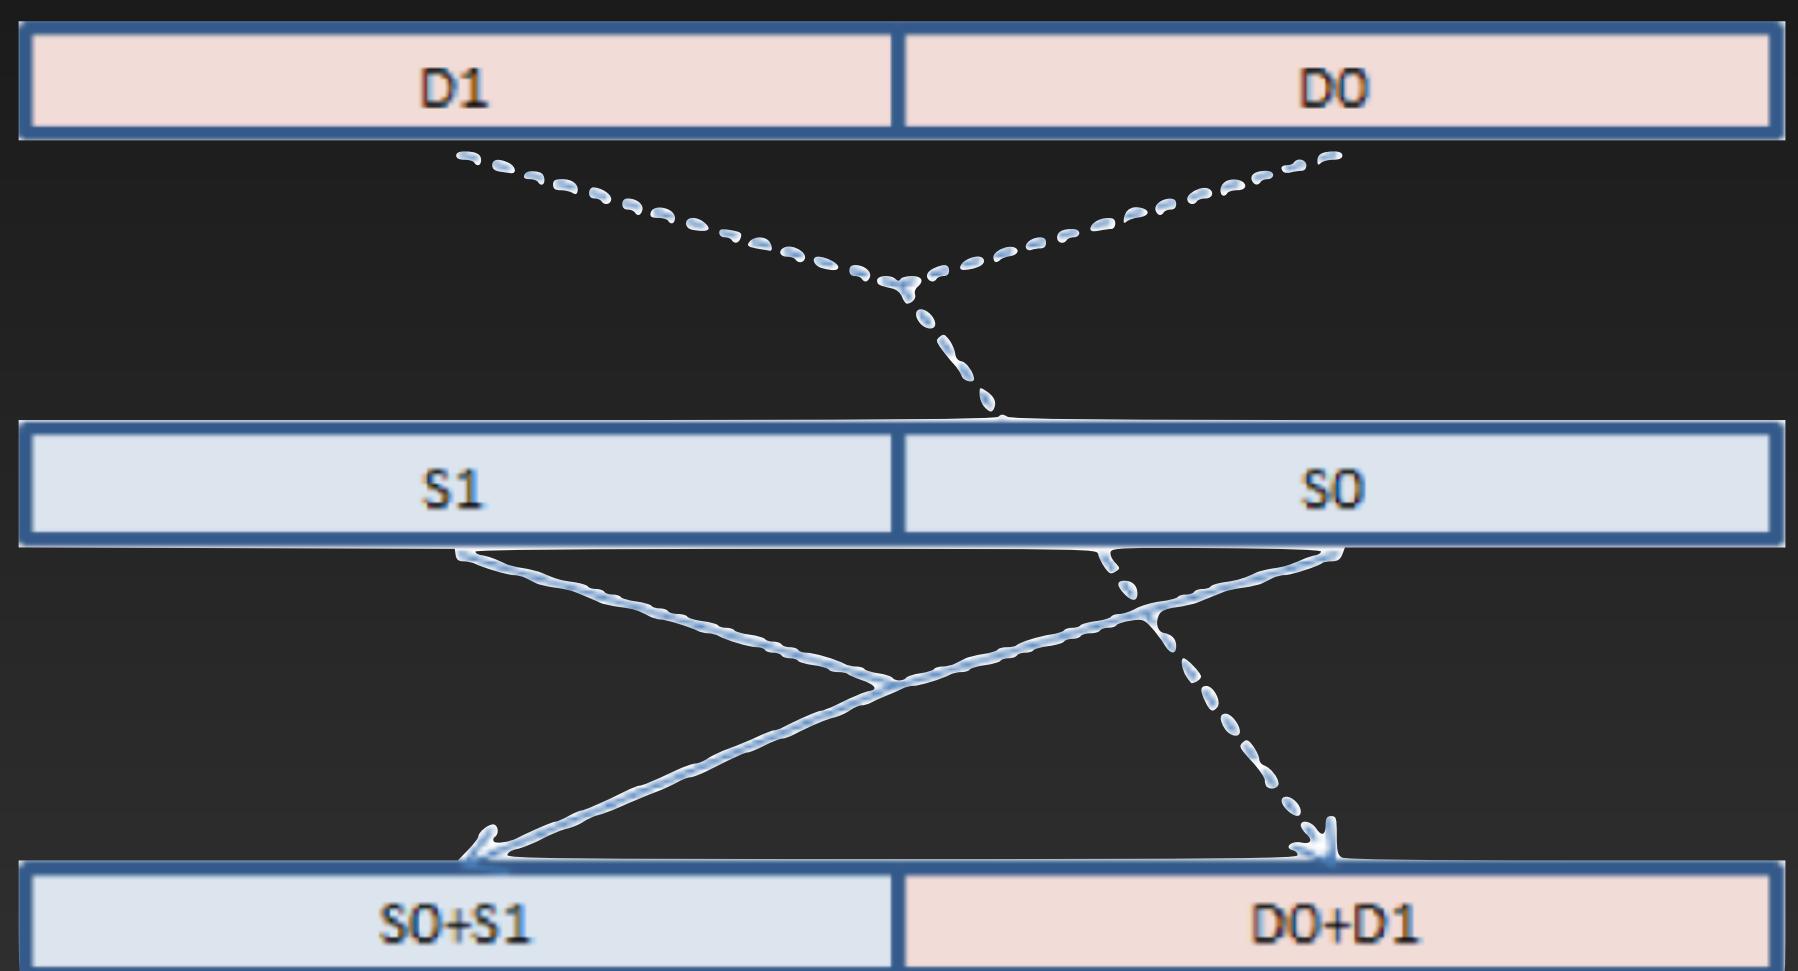
When you, your users, your packages or your libraries don't provide `-march={ foo }`
you are locking your code generation to an ISA from the year 2000

It's hard to believe, but that's over 2 decades ago, and technology has come a long
way

SSE3-2004

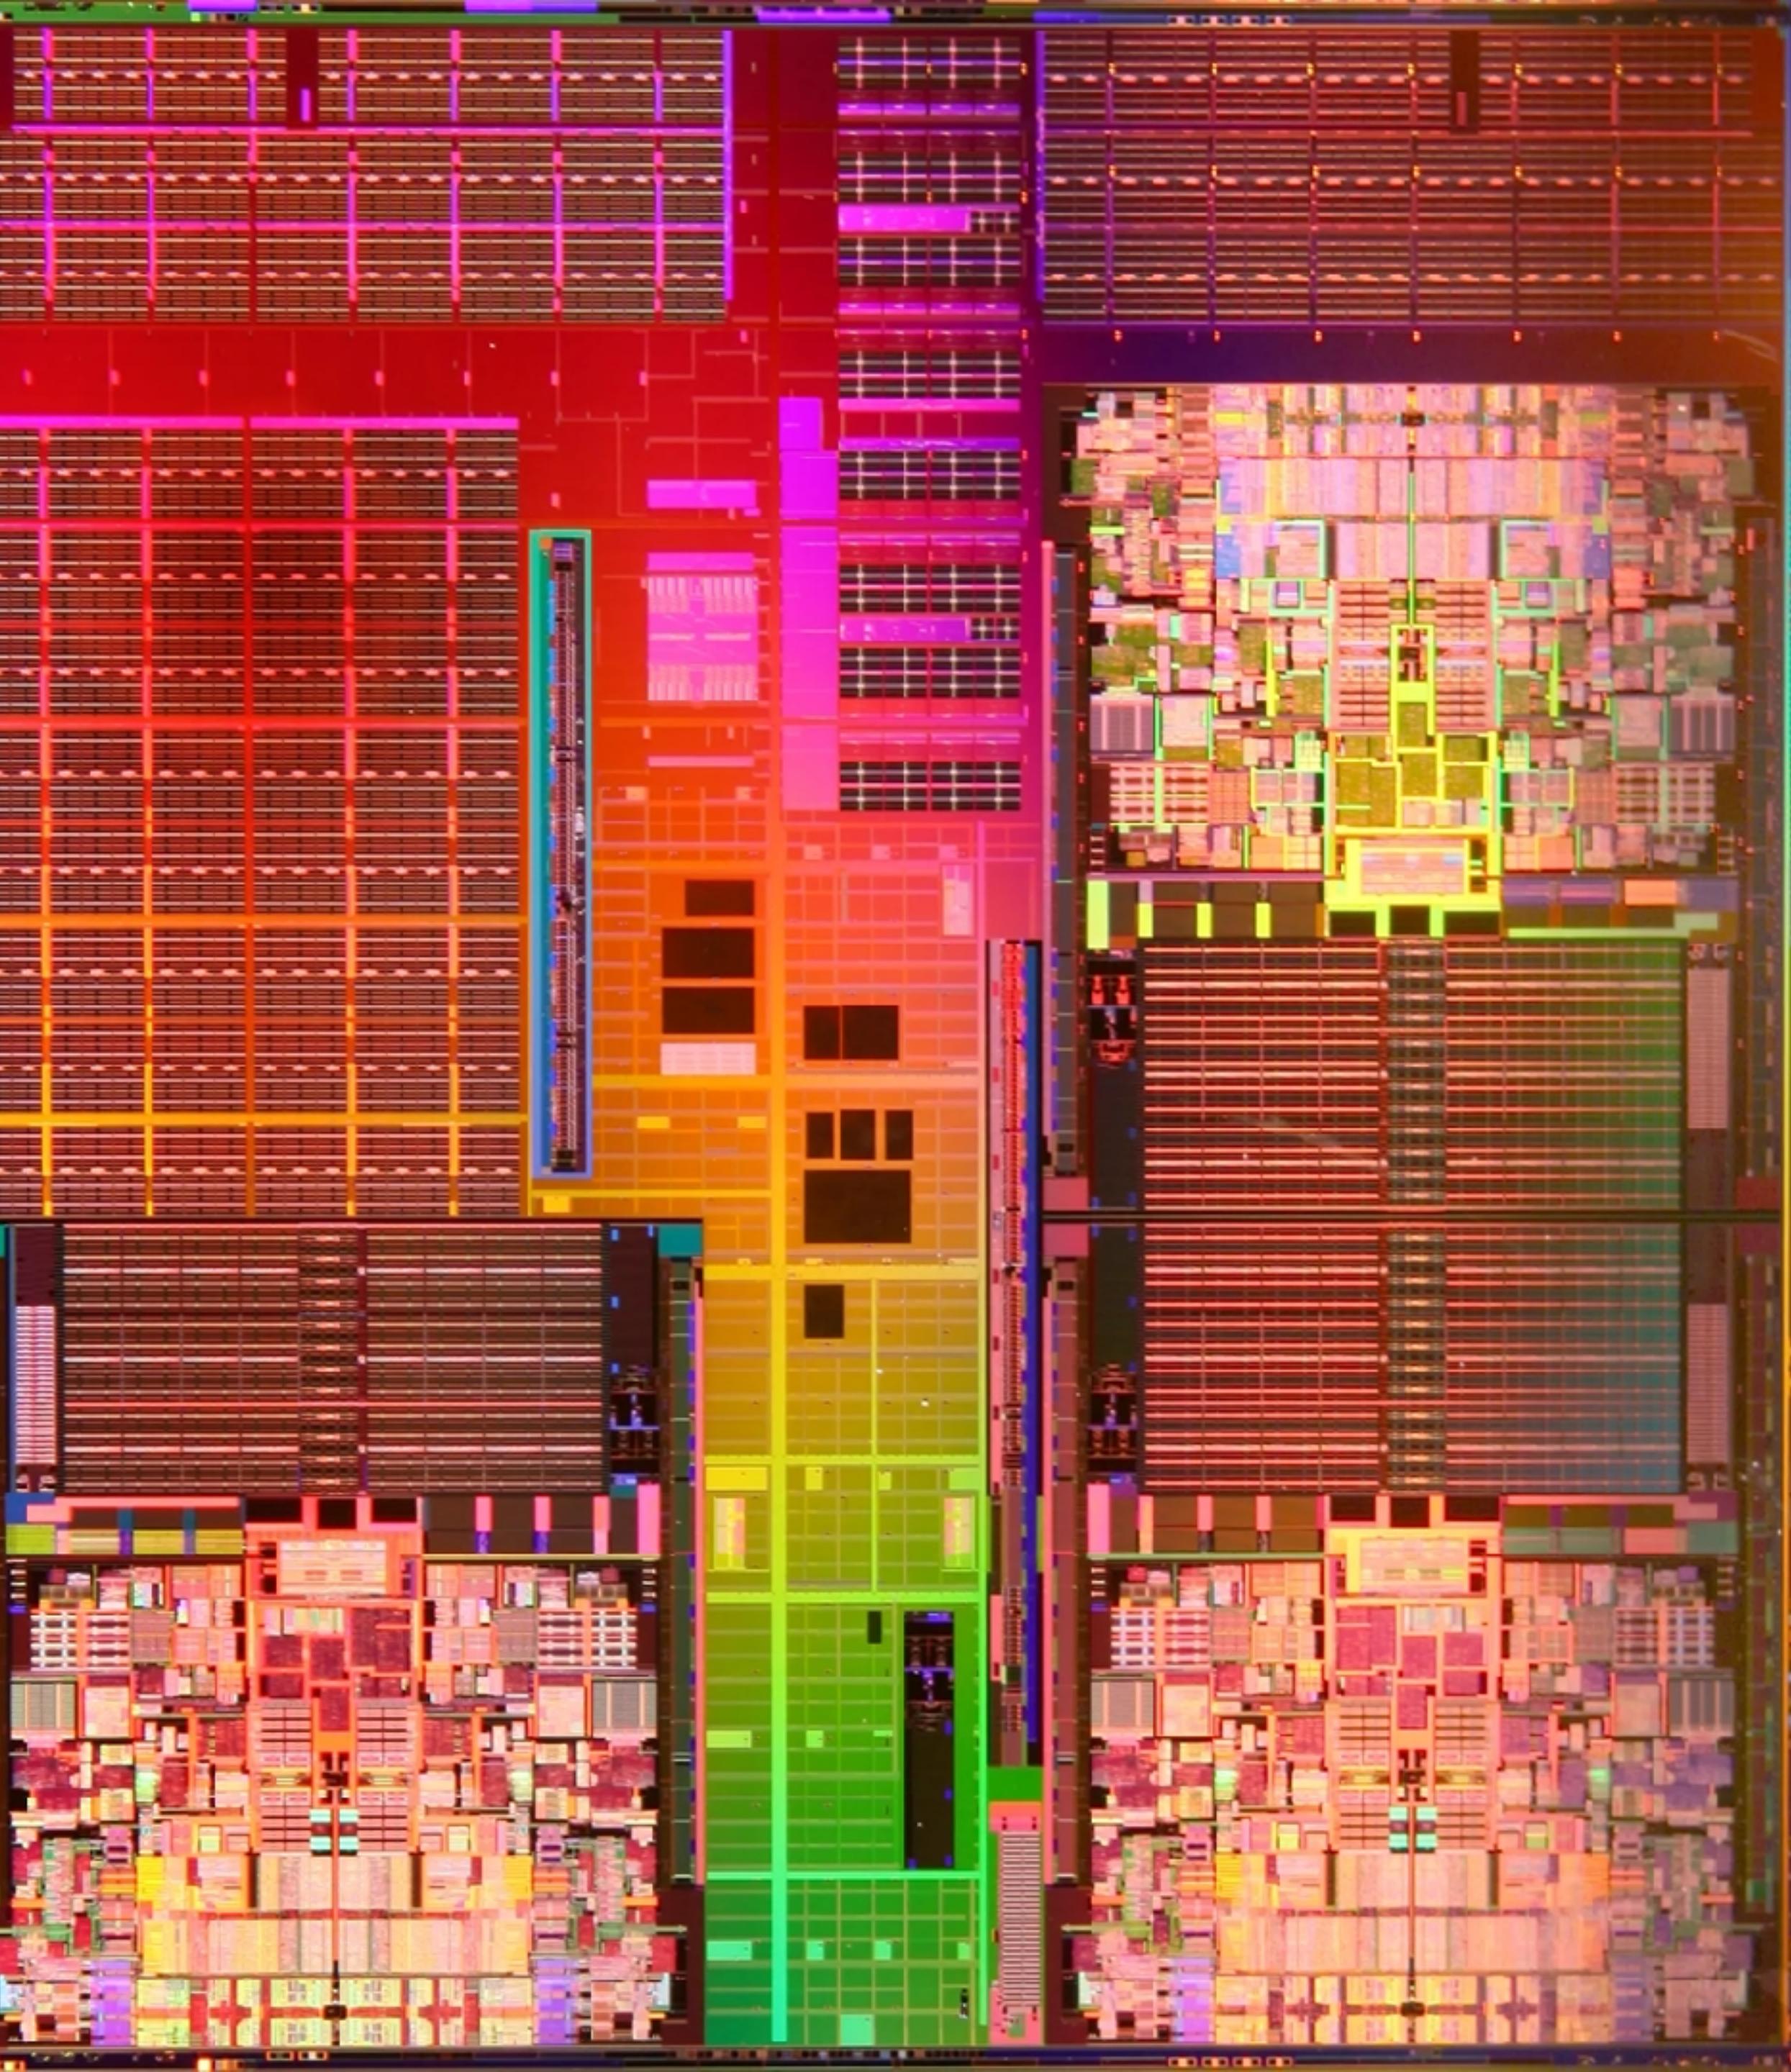
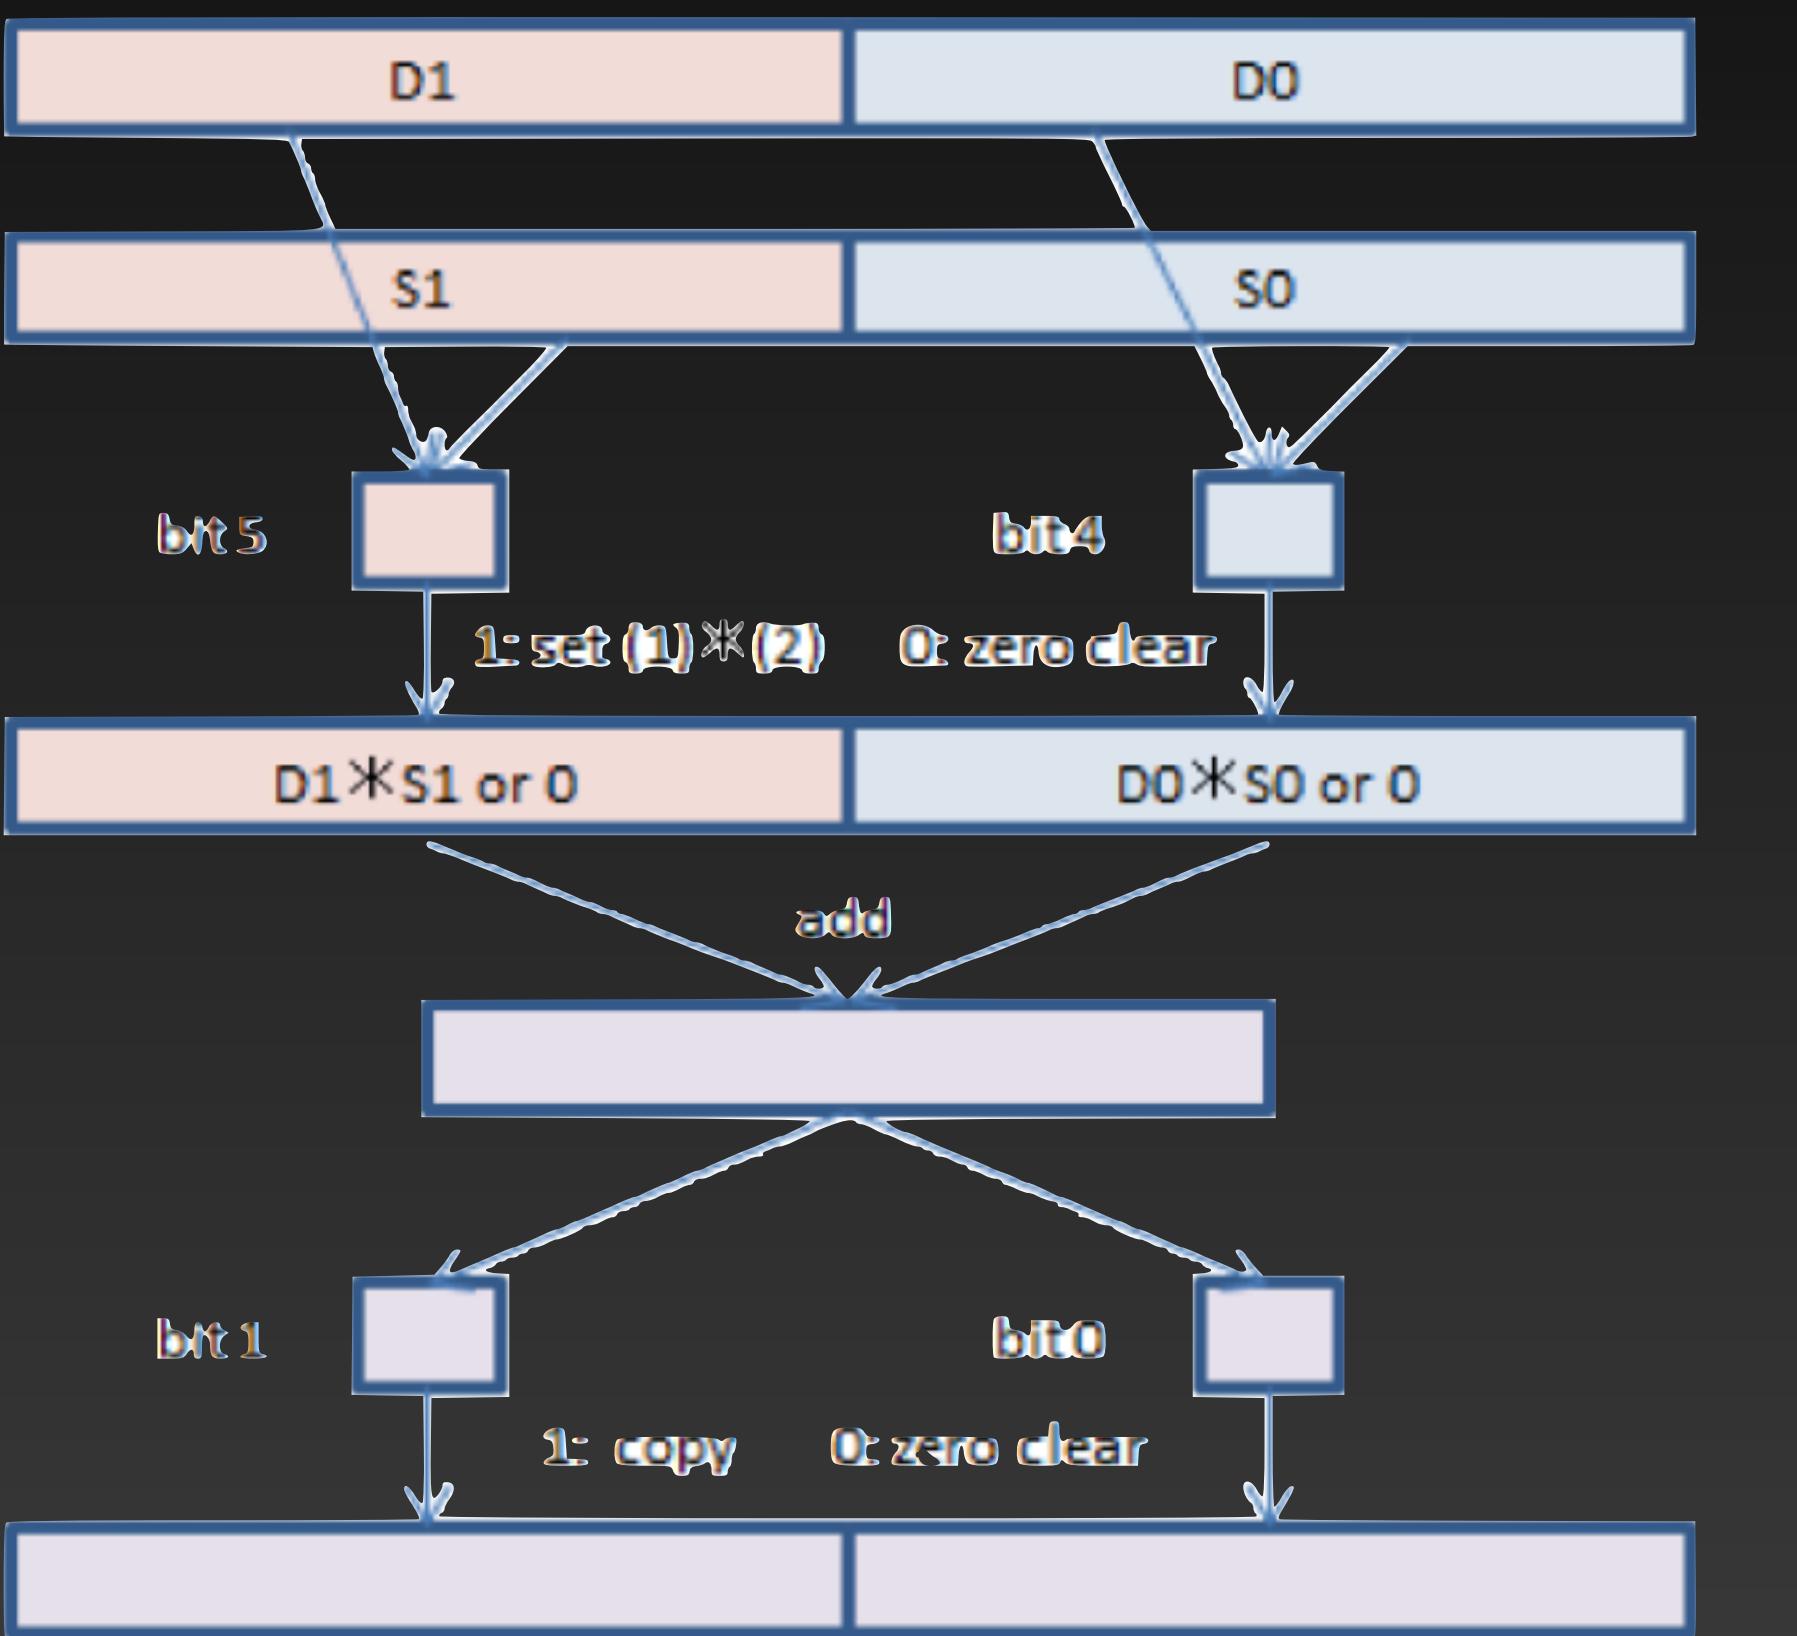
Horizontal sums!

- HADDPD xmm1, xmm2



SSE4.1 and SSE4.2-2004 Dot Products!

- DPPD xmm1, xmm2, imm8



Targets that implement all the way to SSE4.2 are considered x86-64-V2

A “generic” CPU target that includes every CPU, including embed, since ~2013

Provides the compiler with a generic target for ISA’s, doesn’t include any tuning/
implementation data

MISC

“Um actually”- a close friend

- There were also some weird in between extensions
 - AMD 3DNOW!
 - SSSE3
 - SSE5*
 - etc.
- None of them survived to become anything particularly relevant
- With the move to 64 bit operating systems, we double from 8x128 bit SSE registers to 16x128 bit registers registers. You can thank the fine folks at AMD for that! (AMD64)

Section 1+2

Quick clarifications?

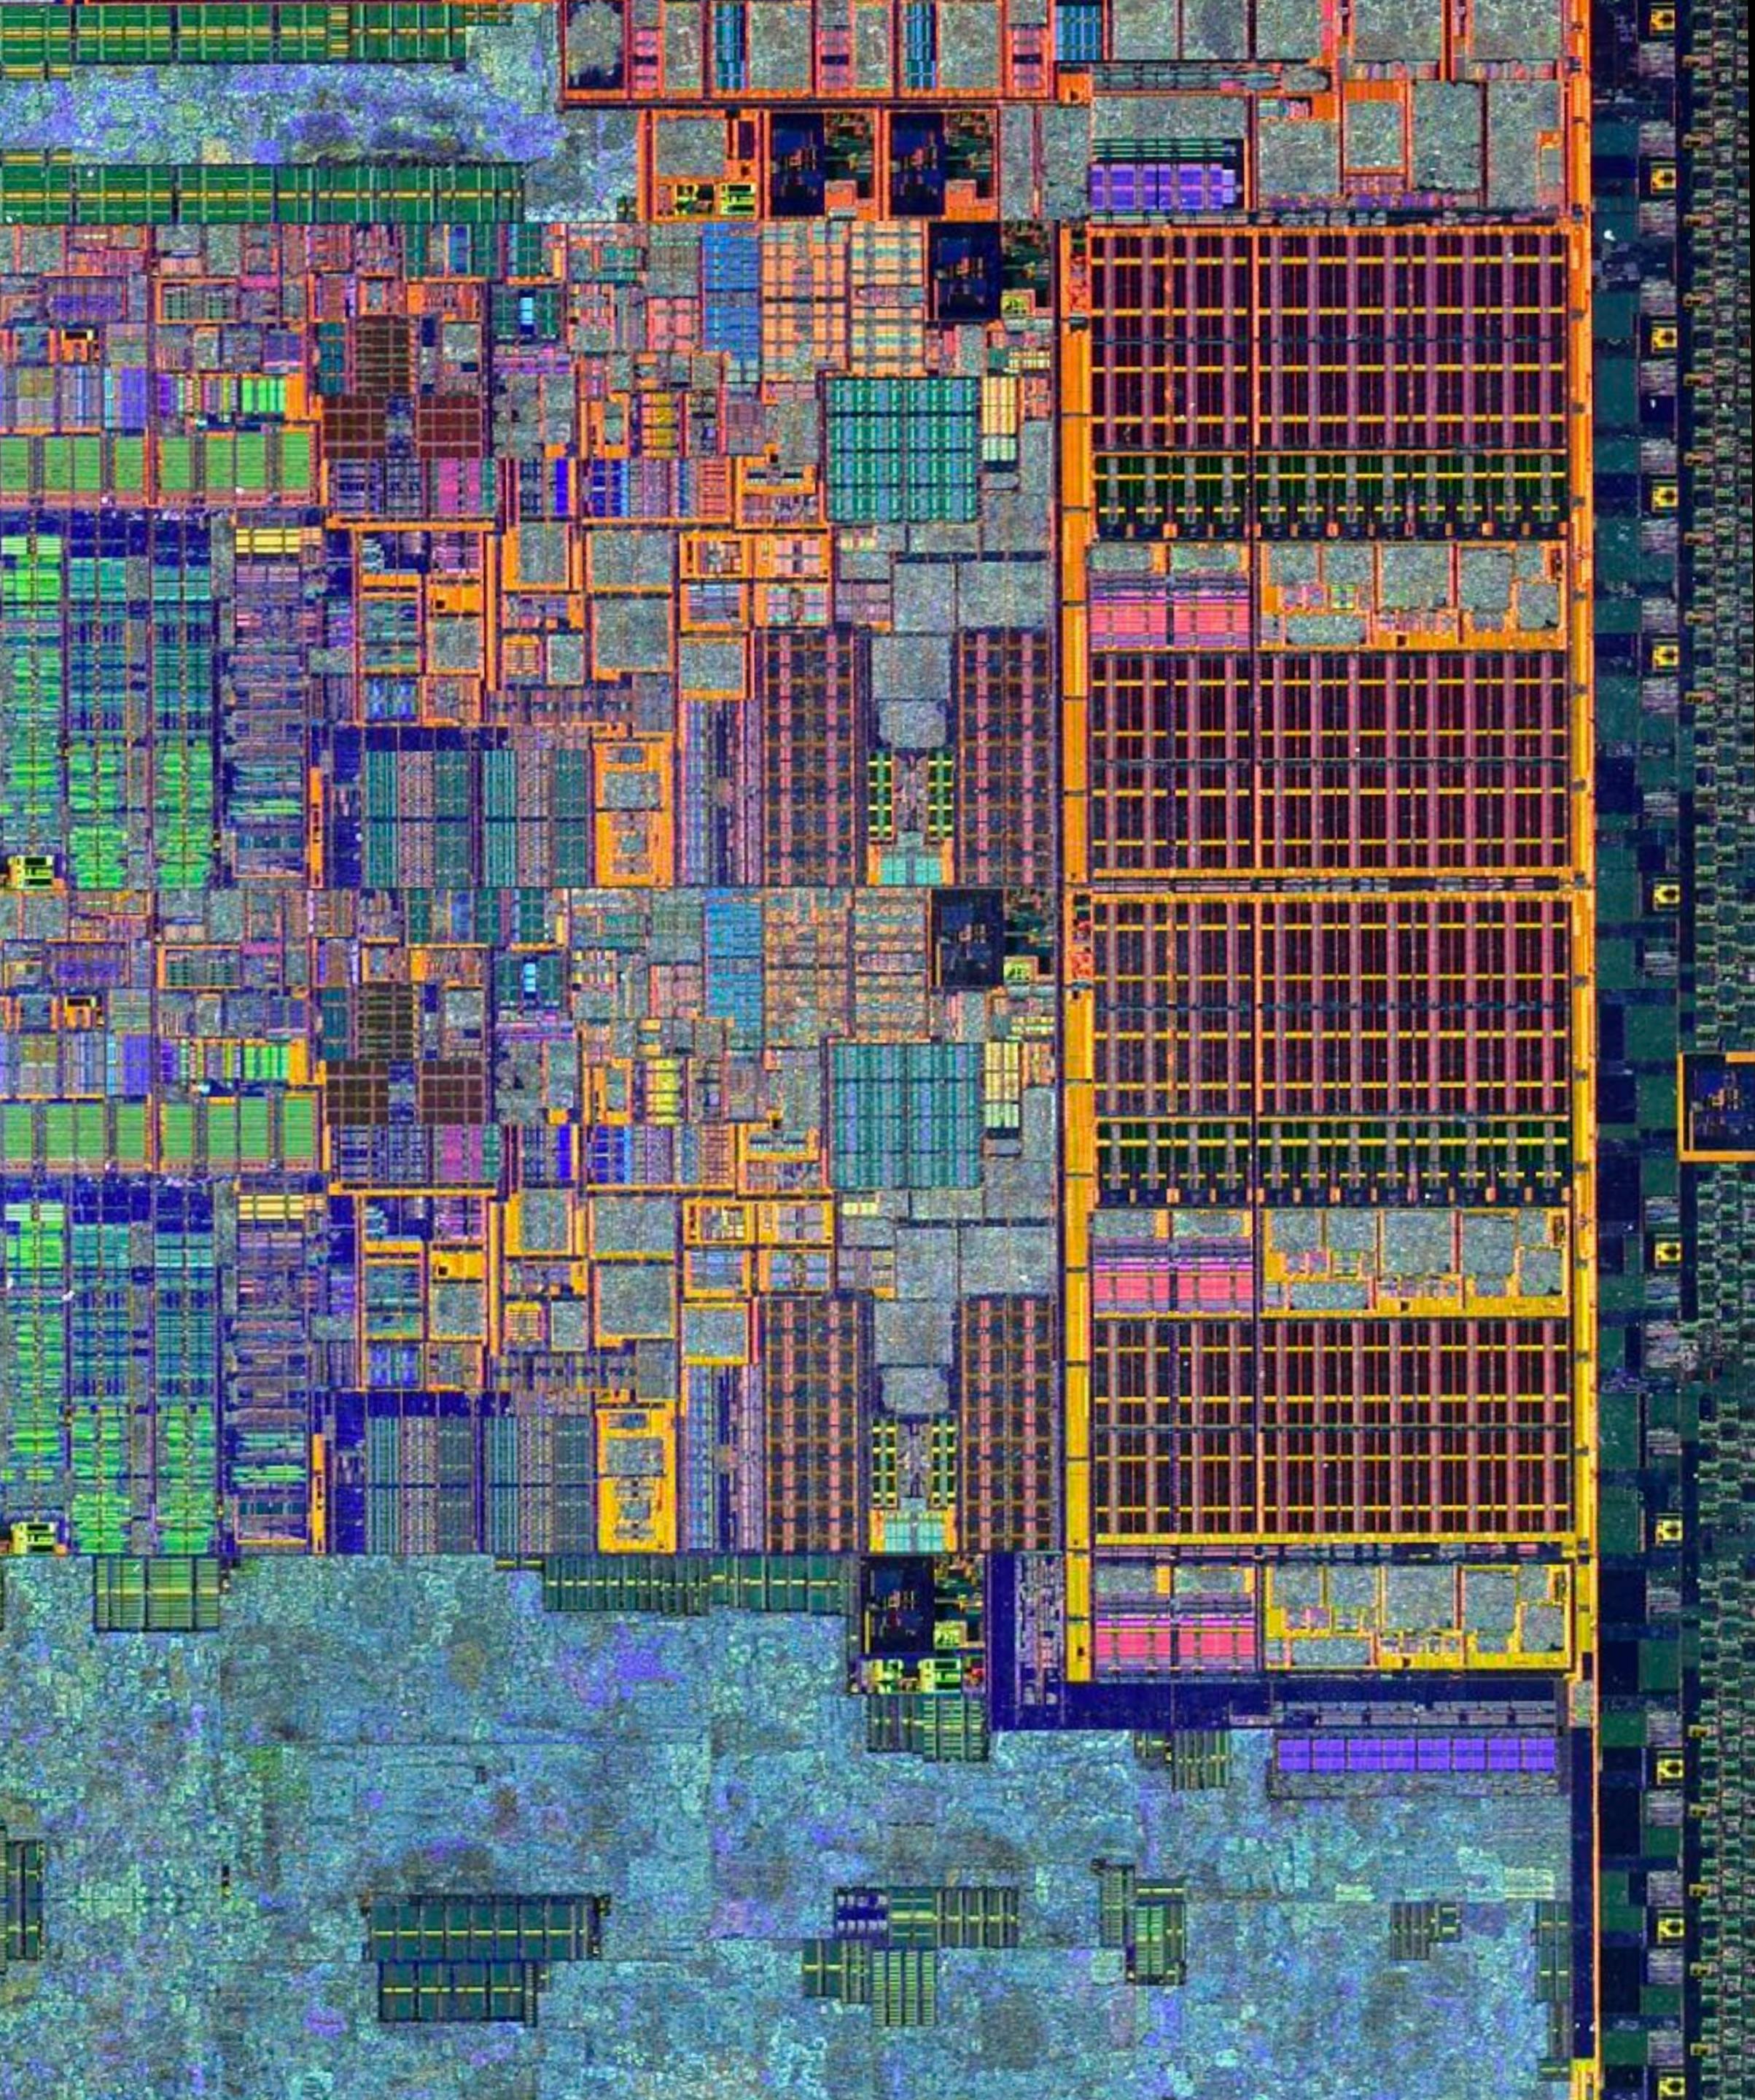
Section 3

Advanced Floating point SIMD

AVX(1) ~2010

“Advanced” SIMD

- 16 x 256 bit vector registers
- 2 data types:
 - 4* FP64
 - 8* FP32
- “Expands” the encoding space with VEX encoding; no more register clobbering!
- addpd -> vaddpd



AVX

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);

}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

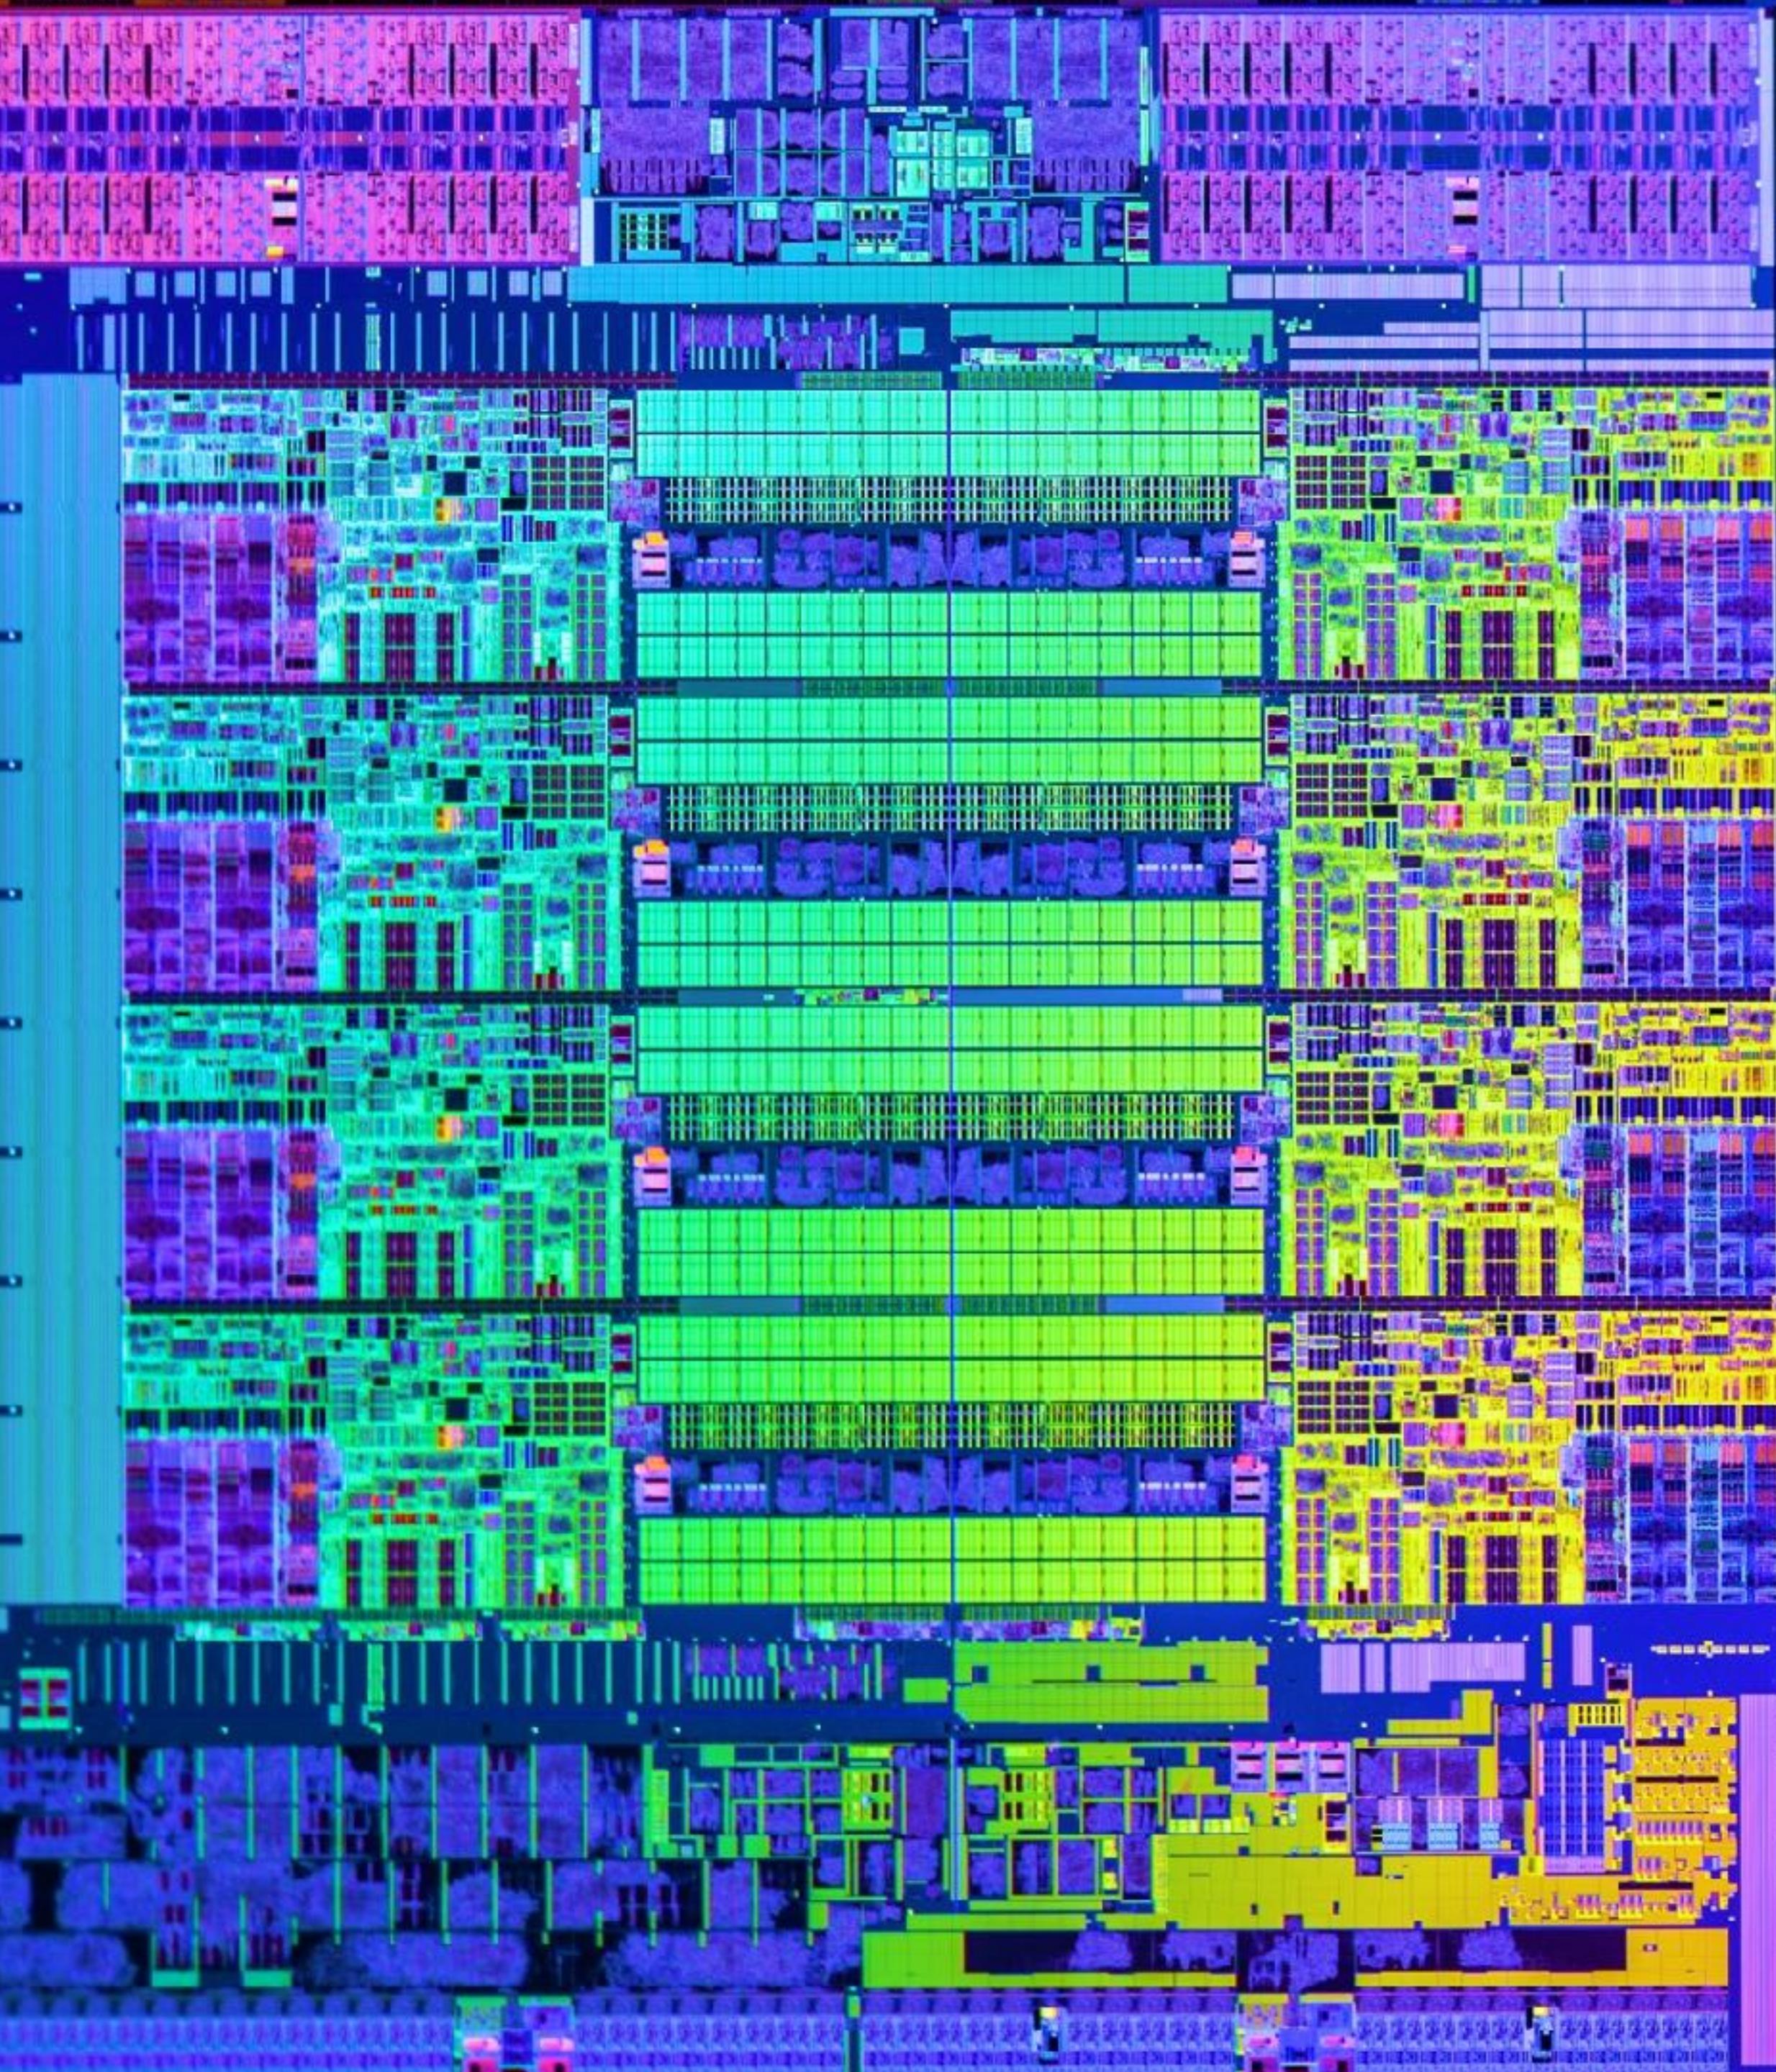
    return 0;
}
```

```
DAXPY(int, double*, double*, double*):
    cmp    edi,0x2
    mov    eax,0x1
    cmovge eax,edi
    cmp    eax,0x10
    jae    17 <DAXPY(int, double*, double*, double*)
+0x17>
    xor    edi,edi
    jmp    de <DAXPY(int, double*, double*, double*)
+0xde>
    lea    rdi,[rcx+rax*8]
    lea    r8,[rsi+rax*8]
    cmp    r8,rcx
    seta   r10b
    lea    r8,[rdx+rax*8]
    cmp    rdi,rsi
    seta   r11b
    cmp    r8,rcx
    seta   r8b
    cmp    rdi,rdx
    seta   r9b
    xor    edi,edi
    test   r10b,r11b
    jne    de <DAXPY(int, double*, double*, double*)
+0xde>
    and    r8b,r9b
    jne    de <DAXPY(int, double*, double*, double*)
+0xde>
    mov    edi,eax
    and    edi,0x7fffffff
    xor    r8d,r8d
    xchg   ax,ax
    vmovupd ymm0,YMMWORD PTR [rdx+r8*8]
    vmovupd ymm1,YMMWORD PTR [rdx+r8*8+0x20]
    vmovupd ymm2,YMMWORD PTR [rdx+r8*8+0x40]
    vmovupd ymm3,YMMWORD PTR [rdx+r8*8+0x60]
    vmulpd ymm0,ymm0,YMMWORD PTR [rcx+r8*8]
    vaddpd ymm0,ymm0,YMMWORD PTR [rsi+r8*8]
    vmulpd ymm1,ymm1,YMMWORD PTR [rcx+r8*8+0x20]
    vaddpd ymm1,ymm1,YMMWORD PTR [rsi+r8*8+0x20]
    vmulpd ymm2,ymm2,YMMWORD PTR [rcx+r8*8+0x40]
    vaddpd ymm2,ymm2,YMMWORD PTR [rsi+r8*8+0x40]
    vmulpd ymm3,ymm3,YMMWORD PTR [rcx+r8*8+0x60]
    vaddpd ymm3,ymm3,YMMWORD PTR [rsi+r8*8+0x60]
    vmovupd YMMWORD PTR [rcx+r8*8],ymm0
    vmovupd YMMWORD PTR [rcx+r8*8+0x20],ymm1
    vmovupd YMMWORD PTR [rcx+r8*8+0x40],ymm2
    vmovupd YMMWORD PTR [rcx+r8*8+0x60],ymm3
    add    r8,0x10
    cmp    rdi,r8
    jne    60 <DAXPY(int, double*, double*, double*)
+0x60>
    cmp    rdi,rax
    je     185 <DAXPY(int, double*, double*, double*)
+0x185>
    mov    r8,rdi
    not    r8
    add    r8,rax
    mov    r9,rax
    and    r9,0x3
    je     10c <DAXPY(int, double*, double*, double*)
+0x10c>
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm0
    inc    rdi
    dec    r9
    jne    f0 <DAXPY(int, double*, double*, double*)
+0xf0>
    cmp    r8,0x3
    jb    185 <DAXPY(int, double*, double*, double*)
+0x185>
    data16 data16 data16 data16 cs nop WORD PTR
[rax+rax*1+0x0]
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x8]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x8]
    vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x10]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x10]
    vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x18]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x18]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x18]
    vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm0
    add    rdi,0x4
    cmp    rax,rdi
    jne    120 <DAXPY(int, double*, double*, double*)
+0x120>
    vzeroupper
    ret
    nop    DWORD PTR [rax+0x0]
main:
    push   r14
    push   rbx
    push   rax
    mov    edi,0x2000
    call   19e <main+0xe>
        R_X86_64_PLT32 operator new[](unsigned
long)-0x4
    mov    rbx,rax
    mov    edi,0x2000
    call   1ab <main+0x1b>
        R_X86_64_PLT32 operator new[](unsigned
long)-0x4
    mov    r14,rax
    mov    edi,0x2000
    call   1b8 <main+0x28>
        R_X86_64_PLT32 operator new[](unsigned
long)-0x4
    mov    edi,0x400
    mov    rsi,rbx
    mov    rdx,r14
    mov    rcx,rax
    call   1cb <main+0x3b>
        R_X86_64_PLT32 DAXPY(int, double*, double*,
double*)-0x4
    xor    eax,eax
    add    rsp,0x8
    pop    rbx
    pop    r14
    ret
```

AVX(2) ~2010

"Advanced" SIMD 2 (electric Boogaloo)

- 256 bit vectors
- 6.5 data types:
 - 4* FP64
 - 8* FP32
 - 4* 64 bit ints
 - 8* 32 bit ints
 - 16* 16 bit ints
 - 32* 8 bit ints
 - Storage and conversion to IEEE FP16
- FMA instructions released separately, allow for single instruction AXPY



AVX2

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);

}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

    return 0;
}
```

```
DAXPY(int, double*, double*, double*):
    cmp    edi,0x2
    mov    eax,0x1
    cmovge eax,edi
    cmp    eax,0x10
    jae    17 <DAXPY(int, double*, double*, double*)
+0x17>
    xor    edi,edi
    jmp    de <DAXPY(int, double*, double*, double*)
+0xde>
    lea    rdi,[rcx+rax*8]
    lea    r8,[rsi+rax*8]
    cmp    r8,rcx
    seta   r10b
    lea    r8,[rdx+rax*8]
    cmp    rdi,rsi
    seta   r11b
    cmp    r8,rcx
    seta   r8b
    cmp    rdi,rdx
    seta   r9b
    xor    edi,edi
    test   r10b,r11b
    jne    de <DAXPY(int, double*, double*, double*)
+0xde>
    and    r8b,r9b
    jne    de <DAXPY(int, double*, double*, double*)
+0xde>
    mov    edi,eax
    and    edi,0x7fffffff
    xor    r8d,r8d
    xchg   ax,ax
    vmovupd ymm0,YMMWORD PTR [rdx+r8*8]
    vmovupd ymm1,YMMWORD PTR [rdx+r8*8+0x20]
    vmovupd ymm2,YMMWORD PTR [rdx+r8*8+0x40]
    vmovupd ymm3,YMMWORD PTR [rdx+r8*8+0x60]
    vmovupd ymm4,YMMWORD PTR [rcx+r8*8]
    vmovupd ymm5,YMMWORD PTR [rcx+r8*8+0x20]
    vmovupd ymm6,YMMWORD PTR [rcx+r8*8+0x40]
    vmovupd ymm7,YMMWORD PTR [rcx+r8*8+0x60]
    vfmadd213pd ymm4,ymm0,YMMWORD PTR [rsi+r8*8]
    vfmadd213pd ymm5,ymm1,YMMWORD PTR [rsi+r8*8+0x20]
    vfmadd213pd ymm6,ymm2,YMMWORD PTR [rsi+r8*8+0x40]
    vfmadd213pd ymm7,ymm3,YMMWORD PTR [rsi+r8*8+0x60]
    vmovupd YMMWORD PTR [rcx+r8*8],ymm4
    vmovupd YMMWORD PTR [rcx+r8*8+0x20],ymm5
    vmovupd YMMWORD PTR [rcx+r8*8+0x40],ymm6
    vmovupd YMMWORD PTR [rcx+r8*8+0x60],ymm7
    add    r8,0x10
    cmp    rdi,r8
    jne    60 <DAXPY(int, double*, double*, double*)
+0x60>
    cmp    rdi,rax
    je     189 <DAXPY(int, double*, double*, double*)
+0x189>
    mov    r8,rdi
    not    r8
    add    r8,rax
    mov    r9,rax
    and    r9,0x3
    je     10d <DAXPY(int, double*, double*, double*)
+0x10d>
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm1
    inc    rdi
    dec    r9
    jne    f0 <DAXPY(int, double*, double*, double*)
+0xf0>
    cmp    r8,0x3
    jb    189 <DAXPY(int, double*, double*, double*)
+0x189>
    data16 data16 data16 cs nop WORD PTR
[rax+rax*1+0x0]
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8]
    vmovsd xmm2,QWORD PTR [rcx+rdi*8+0x8]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm1
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
    vfmadd213sd xmm0,xmm2,QWORD PTR [rsi+rdi*8+0x8]
    vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8+0x10]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x10]
    vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm1
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x18]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8+0x18]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x18]
    vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm1
    add    rdi,0x4
    cmp    rax,rdi
    jne    120 <DAXPY(int, double*, double*, double*)
+0x120>
    vzeroupper
    ret
    nop    DWORD PTR [rax]
main:
    push   r14
    push   rbx
    push   rax
    mov    edi,0x2000
    call   19e <main+0xe>
        R_X86_64_PLT32 operator new[](unsigned
long)-0x4
    mov    rbx,rax
    mov    edi,0x2000
    call   1ab <main+0x1b>
        R_X86_64_PLT32 operator new[](unsigned
long)-0x4
    mov    r14,rax
    mov    edi,0x2000
    call   1b8 <main+0x28>
        R_X86_64_PLT32 operator new[](unsigned
long)-0x4
    mov    edi,0x400
    mov    rsi,rbx
    mov    rdx,r14
    mov    rcx,rax
    call   1cb <main+0x3b>
        R_X86_64_PLT32 DAXPY(int, double*, double*,
double*)-0x4
    xor    eax,eax
    add    rsp,0x8
    pop    rbx
    pop    r14
    ret
```

AVX

Perform DAXPY z[]=(x[]+y[])*z[]

```
DAXPY(int, double*, double*, double*):
    cmp    edi,0x2
    mov    eax,0x1
    cmovge eax,edi
    cmp    eax,0x10
    jae    17 <DAXPY(int, double*, double*, double*)
+0x17>
    xor    edi,edi
    jmp    de <DAXPY(int, double*, double*, double*)
+0xde>
    lea    rdi,[rcx+rax*8]
    lea    r8,[rsi+rax*8]
    cmp    r8,rcx
    seta   r10b
    lea    r8,[rdx+rax*8]
    cmp    rdi,rsi
    seta   r11b
    cmp    r8,rcx
    seta   r8b
    cmp    rdi,rdx
    seta   r9b
    xor    edi,edi
    test   r10b,r11b
    jne    de <DAXPY(int, double*, double*, double*)
+0xde>
    and   r8b,r9b
    jne    de <DAXPY(int, double*, double*, double*)
+0xde>
    mov    edi,eax
    and   edi,0xfffffffff0
    xor    r8d,r8d
    xchg  ax,ax
    vmovupd ymm0,YMMWORD PTR [rdx+r8*8]
    vmovupd ymm1,YMMWORD PTR [rdx+r8*8+0x20]
    vmovupd ymm2,YMMWORD PTR [rdx+r8*8+0x40]
    vmovupd ymm3,YMMWORD PTR [rdx+r8*8+0x60]
    vmulpd ymm0,ymm0,YMMWORD PTR [rcx+r8*8]
    vaddpd ymm0,ymm0,YMMWORD PTR [rsi+r8*8]
    vmulpd ymm1,ymm1,YMMWORD PTR [rcx+r8*8+0x20]
    vaddpd ymm1,ymm1,YMMWORD PTR [rsi+r8*8+0x20]
    vmulpd ymm2,ymm2,YMMWORD PTR [rcx+r8*8+0x40]
    vaddpd ymm2,ymm2,YMMWORD PTR [rsi+r8*8+0x40]
    vmulpd ymm3,ymm3,YMMWORD PTR [rcx+r8*8+0x60]
    vaddpd ymm3,ymm3,YMMWORD PTR [rsi+r8*8+0x60]
    vmovupd YMMWORD PTR [rcx+r8*8],ymm0
    vmovupd YMMWORD PTR [rcx+r8*8+0x20],ymm1
    vmovupd YMMWORD PTR [rcx+r8*8+0x40],ymm2
    vmovupd YMMWORD PTR [rcx+r8*8+0x60],ymm3
    add    r8,0x10
    cmp    rdi,r8
    jne    60 <DAXPY(int, double*, double*, double*)
+0x60>
    cmp    rdi,rax
    je    185 <DAXPY(int, double*, double*, double*)
+0x185>
    mov    r8,rdi
    not   r8
    add    r8,rax
    mov    r9,rax
    and    r9,0x3
    je    10c <DAXPY(int, double*, double*, double*)
+0x10c>
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x8]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x8]
    vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x10]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x10]
    vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x18]
    vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x18]
    vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x18]
    vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm0
    add    rdi,0x4
    cmp    rax,rdi
    jne    120 <DAXPY(int, double*, double*, double*)
+0x120>
    vzeroupper
    ret
    nop    DWORD PTR [rax+0x0]
main:
    push   r14
    push   rbx
    push   rax
    mov    edi,0x2000
    call   19e <main+0xe>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    rbx,rax
    mov    edi,0x2000
    call   1ab <main+0x1b>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    r14,rax
    mov    edi,0x2000
    call   1b8 <main+0x28>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    edi,0x400
    mov    rsi,rbx
    mov    rdx,r14
    mov    rcx,rax
    call   1cb <main+0x3b>
        R_X86_64_PLT32 DAXPY(int, double*, double*, double*)-0x4
    xor    eax,eax
    add    rsp,0x8
    pop    rbx
    pop    r14
    ret

```

AVX2+FMA

Perform DAXPY z[]=(x[]+y[])*z[]

```
DAXPY(int, double*, double*, double*):
    cmp    edi,0x2
    inc    rdi
    dec    r9
    jne    f0 <DAXPY(int, double*, double*, double*)
+0xf0>
    cmp    r8,0x3
    jb    189 <DAXPY(int, double*, double*, double*)
+0x185>
    data16 data16 data16 cs nop WORD PTR [rax+rax*1+0x0]
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8]
    vmovsd xmm2,QWORD PTR [rcx+rdi*8+0x8]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm1
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
    vfmadd213sd xmm0,xmm2,QWORD PTR [rsi+rdi*8+0x8]
    vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
    vfmadd213sd xmm0,xmm1,QWORD PTR [rsi+rdi*8+0x10]
    vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x18]
    vfmadd213sd xmm0,xmm2,QWORD PTR [rsi+rdi*8+0x18]
    vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm0
    add    rdi,0x4
    cmp    rax,rdi
    jne    120 <DAXPY(int, double*, double*, double*)
+0x120>
    vzeroupper
    ret
    nop    DWORD PTR [rax]
main:
    push   r14
    push   rbx
    push   rax
    mov    edi,0x2000
    call   19e <main+0xe>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    rbx,rax
    mov    edi,0x2000
    call   1ab <main+0x1b>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    r14,rax
    mov    edi,0x2000
    call   1b8 <main+0x28>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    edi,0x400
    mov    rsi,rbx
    mov    rdx,r14
    mov    rcx,rax
    call   1cb <main+0x3b>
        R_X86_64_PLT32 DAXPY(int, double*, double*, double*)-0x4
    xor    eax,eax
    add    rsp,0x8
    pop    rbx
    pop    r14
    ret

```

AVX

Perform DAXPY $z[] = (x[] + y[]) * z[]$

Main Loop:

```

vmovupd ymm0,YMMWORD PTR [rdx+r8*8]
vmovupd ymm1,YMMWORD PTR [rdx+r8*8+0x20]
vmovupd ymm2,YMMWORD PTR [rdx+r8*8+0x40]
vmovupd ymm3,YMMWORD PTR [rdx+r8*8+0x60]
vmulpd ymm0,ymm0,YMMWORD PTR [rcx+r8*8]
vaddpd ymm0,ymm0,YMMWORD PTR [rsi+r8*8]
vmulpd ymm1,ymm1,YMMWORD PTR [rcx+r8*8+0x20]
vaddpd ymm1,ymm1,YMMWORD PTR [rsi+r8*8+0x20]
vmulpd ymm2,ymm2,YMMWORD PTR [rcx+r8*8+0x40]
vaddpd ymm2,ymm2,YMMWORD PTR [rsi+r8*8+0x40]
vmulpd ymm3,ymm3,YMMWORD PTR [rcx+r8*8+0x60]
vaddpd ymm3,ymm3,YMMWORD PTR [rsi+r8*8+0x60]
vmovupd YMMWORD PTR [rcx+r8*8],ymm0
vmovupd YMMWORD PTR [rcx+r8*8+0x20],ymm1
vmovupd YMMWORD PTR [rcx+r8*8+0x40],ymm2
vmovupd YMMWORD PTR [rcx+r8*8+0x60],ymm3
add    r8,0x10

```

Residuals:

```

vmovsd xmm0,QWORD PTR [rdx+rdi*8]
vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8]
vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8]
vmovsd QWORD PTR [rcx+rdi*8],xmm0
vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x8]
vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x8]
vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x10]
vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x10]
vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm0
vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x18]
vmulsd xmm0,xmm0,QWORD PTR [rcx+rdi*8+0x18]
vaddsd xmm0,xmm0,QWORD PTR [rsi+rdi*8+0x18]
vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm0
add    rdi,0x4
cmp    rax,rdi
jne    120 <DAXPY(int, double*, double*, double*)+0x120>
vzeroupper
ret

```

AVX2+FMA

Perform DAXPY $z[] = (x[] + y[]) * z[]$

Main Loop:

```

vmovupd ymm0,YMMWORD PTR [rdx+r8*8]
vmovupd ymm1,YMMWORD PTR [rdx+r8*8+0x20]
vmovupd ymm2,YMMWORD PTR [rdx+r8*8+0x40]
vmovupd ymm3,YMMWORD PTR [rdx+r8*8+0x60]
vmovupd ymm4,YMMWORD PTR [rcx+r8*8]
vmovupd ymm5,YMMWORD PTR [rcx+r8*8+0x20]
vmovupd ymm6,YMMWORD PTR [rcx+r8*8+0x40]
vmovupd ymm7,YMMWORD PTR [rcx+r8*8+0x60]
vfmaadd213pd ymm4,ymm0,YMMWORD PTR [rsi+r8*8]
vfmaadd213pd ymm5,ymm1,YMMWORD PTR [rsi+r8*8+0x20]
vfmaadd213pd ymm6,ymm2,YMMWORD PTR [rsi+r8*8+0x40]
vfmaadd213pd ymm7,ymm3,YMMWORD PTR [rsi+r8*8+0x60]
vmovupd YMMWORD PTR [rcx+r8*8],ymm4
vmovupd YMMWORD PTR [rcx+r8*8+0x20],ymm5
vmovupd YMMWORD PTR [rcx+r8*8+0x40],ymm6
vmovupd YMMWORD PTR [rcx+r8*8+0x60],ymm7

```

Residuals:

```

vmovsd xmm0,QWORD PTR [rdx+rdi*8]
vmovsd xmm1,QWORD PTR [rcx+rdi*8]
vmovsd xmm2,QWORD PTR [rcx+rdi*8+0x8]
vfmaadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8]
vmovsd QWORD PTR [rcx+rdi*8],xmm1
vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
vfmaadd213sd xmm0,xmm2,QWORD PTR [rsi+rdi*8+0x8]
vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
vmovsd xmm1,QWORD PTR [rcx+rdi*8+0x10]
vmovsd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x10]
vfmaadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x10]
vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm1
vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
vmovsd xmm1,QWORD PTR [rcx+rdi*8+0x18]
vmovsd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x18]
vfmaadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x18]
vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm1
add    rdi,0x4
cmp    rax,rdi
jne    120 <DAXPY(int, double*, double*, double*)+0x120>
vzeroupper
ret

```

Targets that implement all the way to AVX2 including FMA3 are considered x86-64-V3 compliant

A “generic” CPU target that includes every major CPU since ~2014

Provides the compiler with a generic target for ISA's, doesn't include any tuning/implementation data

Section 3

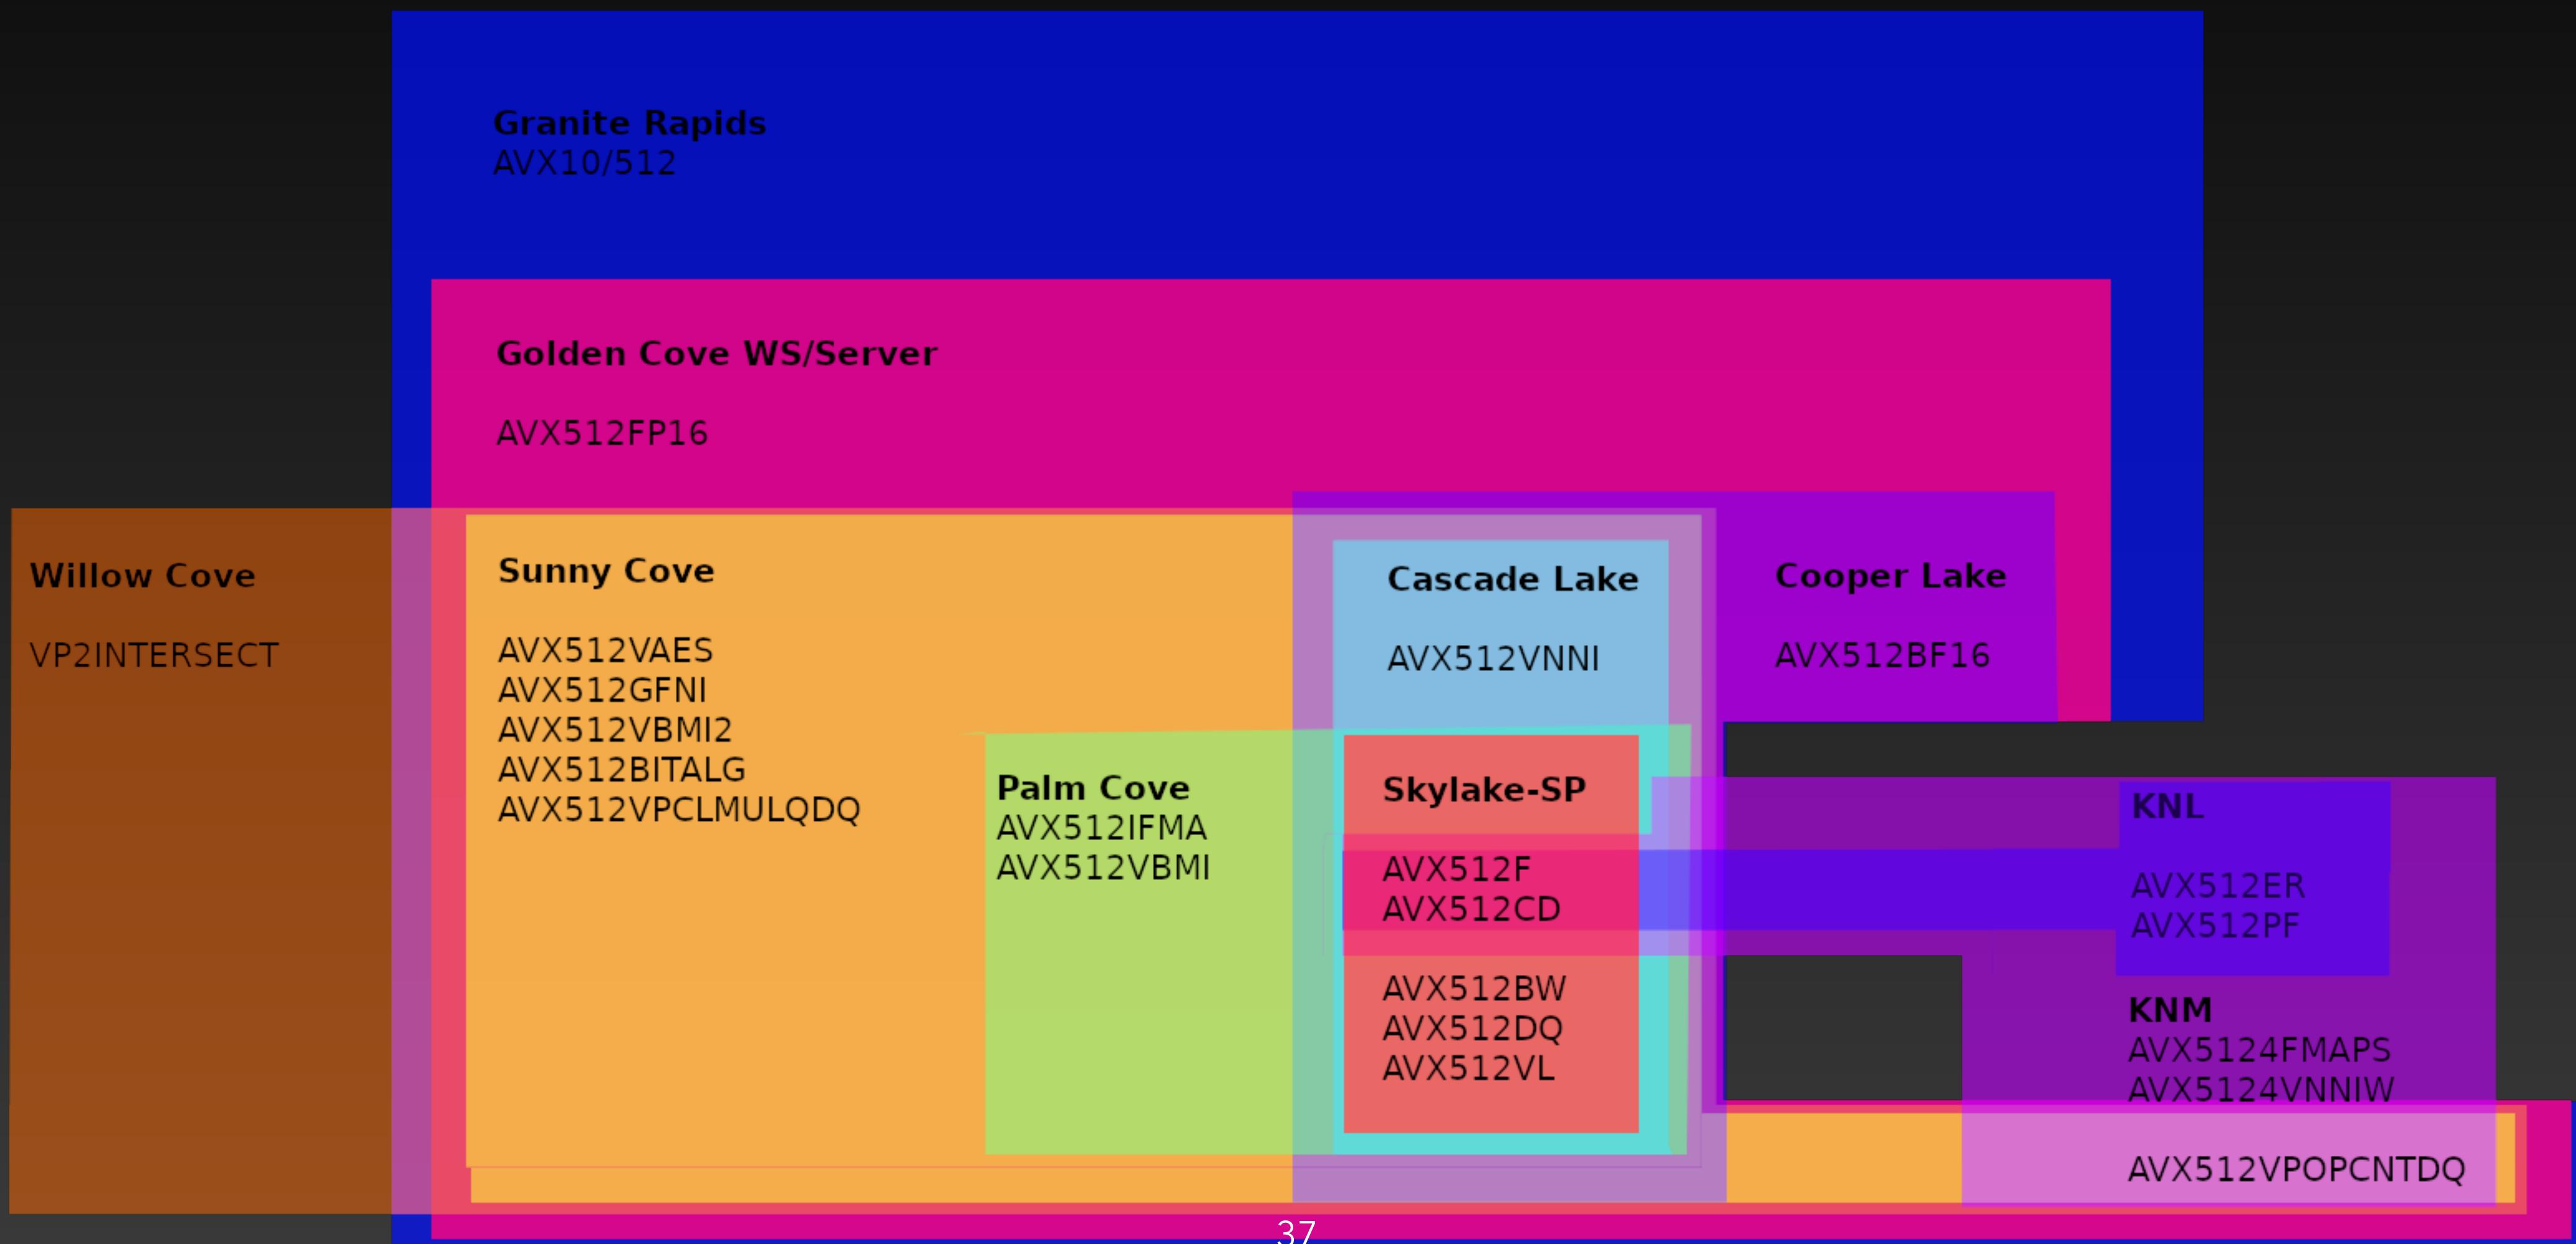
Quick clarifications?

Section 4

Larger Advanced FP SIMD

Section 4

AVX “3” to “AVX9”



I'm ignoring Phi

For the sake of this talk, Phi and Larrabee aren't super relevant

If you're interested in what became Phi, I highly recommend Tom Forsyths talk
“SMACNI to AVX512: The Life cycle of an instruction set”.
Links are one the EasyBuild page!

AVX512 AKA AVX"3"

"Advanced" SIMD 3 (Big vectors, look at me!)

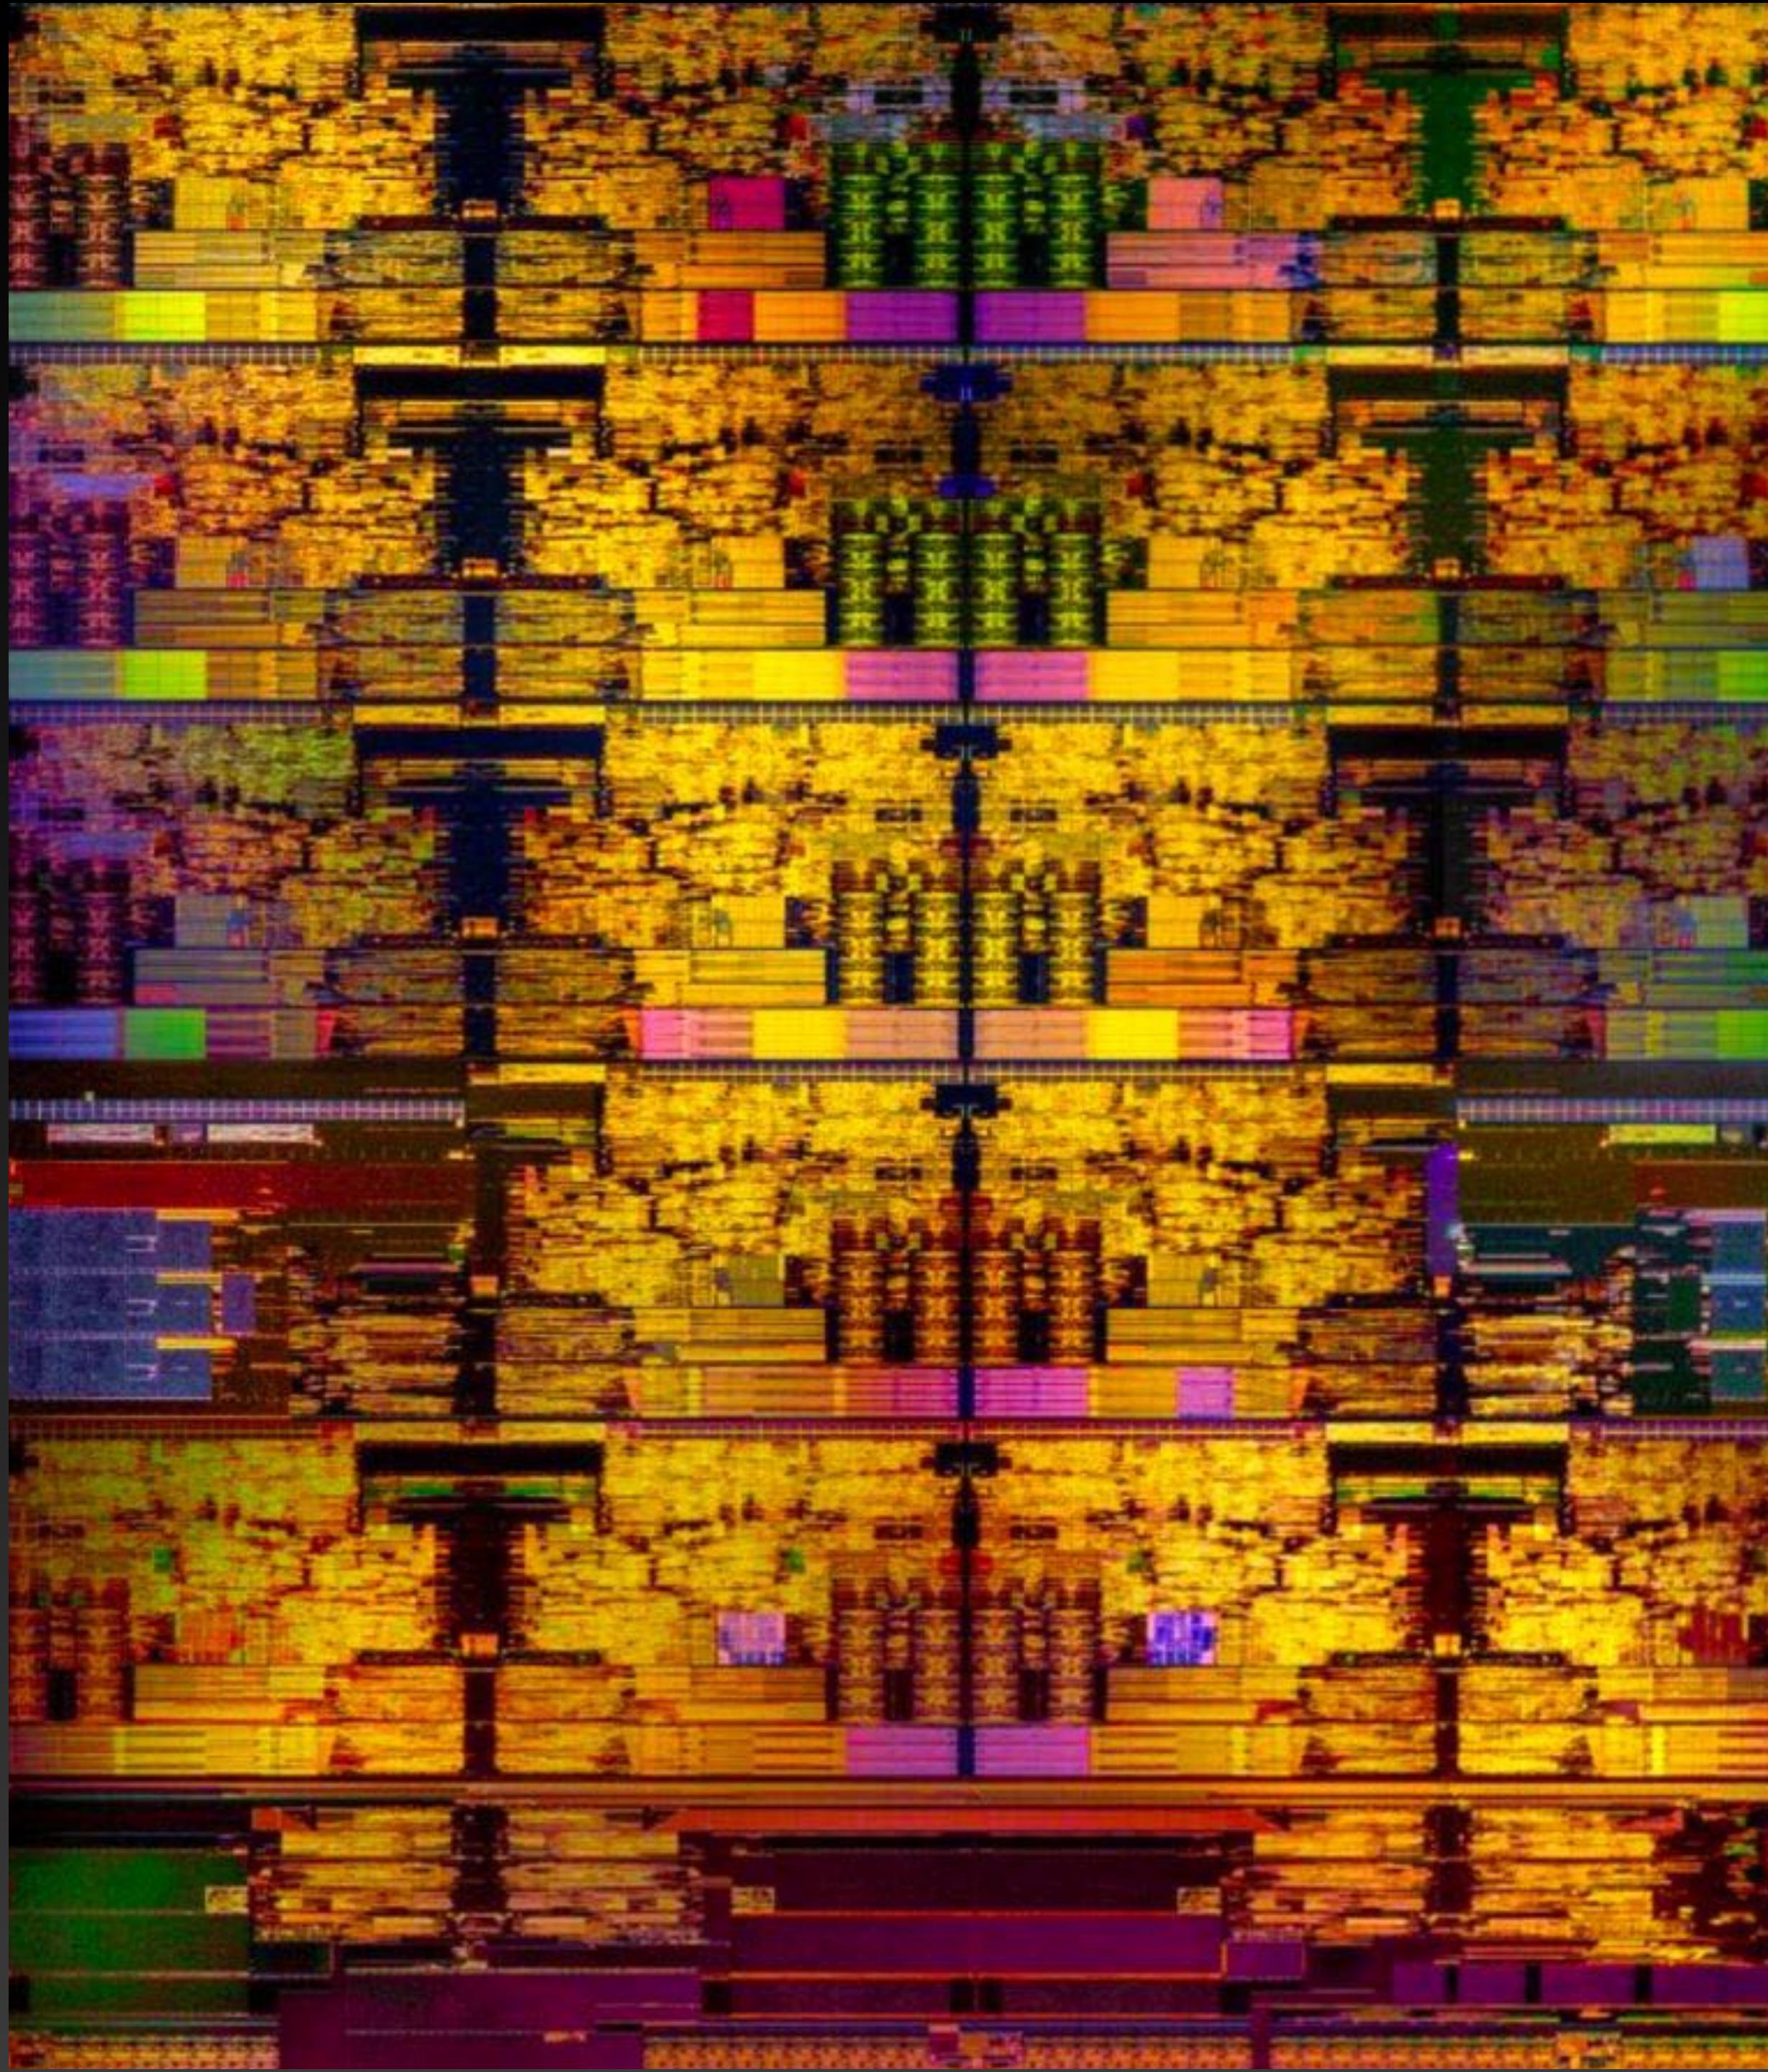
- 32x512 bit "ZMM" vector registers
- 4 data types:
 - 8* FP64
 - 16* FP32
 - 8* 64 bit ints
 - 16* 32 bit ints



“Real” AVX512 (V4)

Variable length, Bytes/words/DWords/QWords

- 5 data types:
 - 16* 32 bit uints
 - 32* 16 bit (u)ints
 - 64 * 8 bit (u)ints
- HPC: more of the same from AVX1



AVX “4”

Perform DAXPY $z[] = (x[] + y[]) * z[]$

```
// function to add the elements of two arrays
void DAXPY(int n, double *x, double *y, double *z)
{
    int i = 0;
    do {z[i] = (x[i] + y[i]) * z[i];
        i++;} while(i<n);

}

int main(void)
{
    int N = 1<<10; // ~1000 elements

    double *x = new double[N];
    double *y = new double[N];
    double *z = new double[N];

    // Run kernel on 1K elements
    DAXPY(N, x, y, z);

    return 0;
}
```

```
DAXPY(int, double*, double*, double*):
    cmp    edi,0x2
    mov    eax,0x1
    cmovge eax,edi
    cmp    eax,0x20
    jae    17 <DAXPY(int, double*, double*, double*)
+0x17>
    xor    edi,edi
    jmp    f2 <DAXPY(int, double*, double*, double*)
+0xf2>
    lea    rdi,[rcx+rax*8]
    lea    r8,[rsi+rax*8]
    lea    r9,[rdx+rax*8]
    cmp    r8,rcx
    seta   r10b
    cmp    rdi,rsi
    seta   r11b
    cmp    r9,rcx
    seta   r8b
    cmp    rdi,rdx
    seta   r9b
    xor    edi,edi
    test   r10b,r11b
    jne    f2 <DAXPY(int, double*, double*, double*)
+0xf2>
    and    r8b,r9b
    jne    f2 <DAXPY(int, double*, double*, double*)
+0xf2>
    mov    edi,eax
    and    edi,0x7fffffff0
    xor    r8d,r8d
    xchg   ax,ax
    vmovupd zmm0,ZMMWORD PTR [rdx+r8*8]
    vmovupd zmm1,ZMMWORD PTR [rdx+r8*8+0x40]
    vmovupd zmm2,ZMMWORD PTR [rdx+r8*8+0x80]
    vmovupd zmm3,ZMMWORD PTR [rdx+r8*8+0xc0]
    vmovupd zmm4,ZMMWORD PTR [rcx+r8*8]
    vmovupd zmm5,ZMMWORD PTR [rcx+r8*8+0x40]
    vmovupd zmm6,ZMMWORD PTR [rcx+r8*8+0x80]
    vmovupd zmm7,ZMMWORD PTR [rcx+r8*8+0xc0]
    vfmadd213pd zmm4,zmm0,ZMMWORD PTR [rsi+r8*8]
    vfmadd213pd zmm5,zmm1,ZMMWORD PTR [rsi+r8*8+0x40]
    vfmadd213pd zmm6,zmm2,ZMMWORD PTR [rsi+r8*8+0x80]
    vfmadd213pd zmm7,zmm3,ZMMWORD PTR [rsi+r8*8+0xc0]
    vmovupd ZMMWORD PTR [rcx+r8*8],zmm4
    vmovupd ZMMWORD PTR [rcx+r8*8+0x40],zmm5
    vmovupd ZMMWORD PTR [rcx+r8*8+0x80],zmm6
    vmovupd ZMMWORD PTR [rcx+r8*8+0xc0],zmm7
    add    r8,0x20
    cmp    rdi,r8
    jne    60 <DAXPY(int, double*, double*, double*)
+0x60>
    cmp    rdi,rax
    je     1a9 <DAXPY(int, double*, double*, double*)
+0xa9>
    mov    r8,rdi
    not    r8
    add    r8,rax
    mov    r9,rax
    and    r9,0x3
    je     12d <DAXPY(int, double*, double*, double*)
+0x12d>
    data16 cs nop WORD PTR [rax+rax*1+0x0]
    vmovsd xmm0,QWORD PTR [rdx+rdi*8]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8]
    vmovsd QWORD PTR [rcx+rdi*8],xmm1
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x8]
    vfmadd213sd xmm0,xmm2,QWORD PTR [rsi+rdi*8+0x8]
    vmovsd QWORD PTR [rcx+rdi*8+0x8],xmm0
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x10]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8+0x10]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x10]
    vmovsd QWORD PTR [rcx+rdi*8+0x10],xmm1
    vmovsd xmm0,QWORD PTR [rdx+rdi*8+0x18]
    vmovsd xmm1,QWORD PTR [rcx+rdi*8+0x18]
    vfmadd213sd xmm1,xmm0,QWORD PTR [rsi+rdi*8+0x18]
    vmovsd QWORD PTR [rcx+rdi*8+0x18],xmm1
    add    rdi,0x4
    cmp    rax,rdi
    jne    140 <DAXPY(int, double*, double*, double*)
+0x140>
    vzeroupper
    ret
    nop    DWORD PTR [rax]
main:
    push   r14
    push   rbx
    push   rax
    mov    edi,0x800000
    call   1be <main+0xe>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    rbx,rax
    mov    edi,0x800000
    call   1cb <main+0x1b>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    r14,rax
    mov    edi,0x800000
    call   1d8 <main+0x28>
        R_X86_64_PLT32 operator new[](unsigned long)-0x4
    mov    edi,0x100000
    mov    rsi,rbx
    mov    rdx,r14
    mov    rcx,rax
    call   1eb <main+0x3b>
        R_X86_64_PLT32 DAXPY(int, double*, double*, double*)-0x4
    xor    eax,eax
    add    rsp,0x8
    pop    rbx
    pop    r14
    ret
```

AVX512 “V5”

Integer FMA+VBMI

- 1 data type:
 - 8*52 bit integers with 104 bit internal rounding
 - Think of an integer that uses the FP64 unit; codes that can use integer maths love this one

No fun pictures, Canon lake never “really” made it to market

AVX512 “V6”

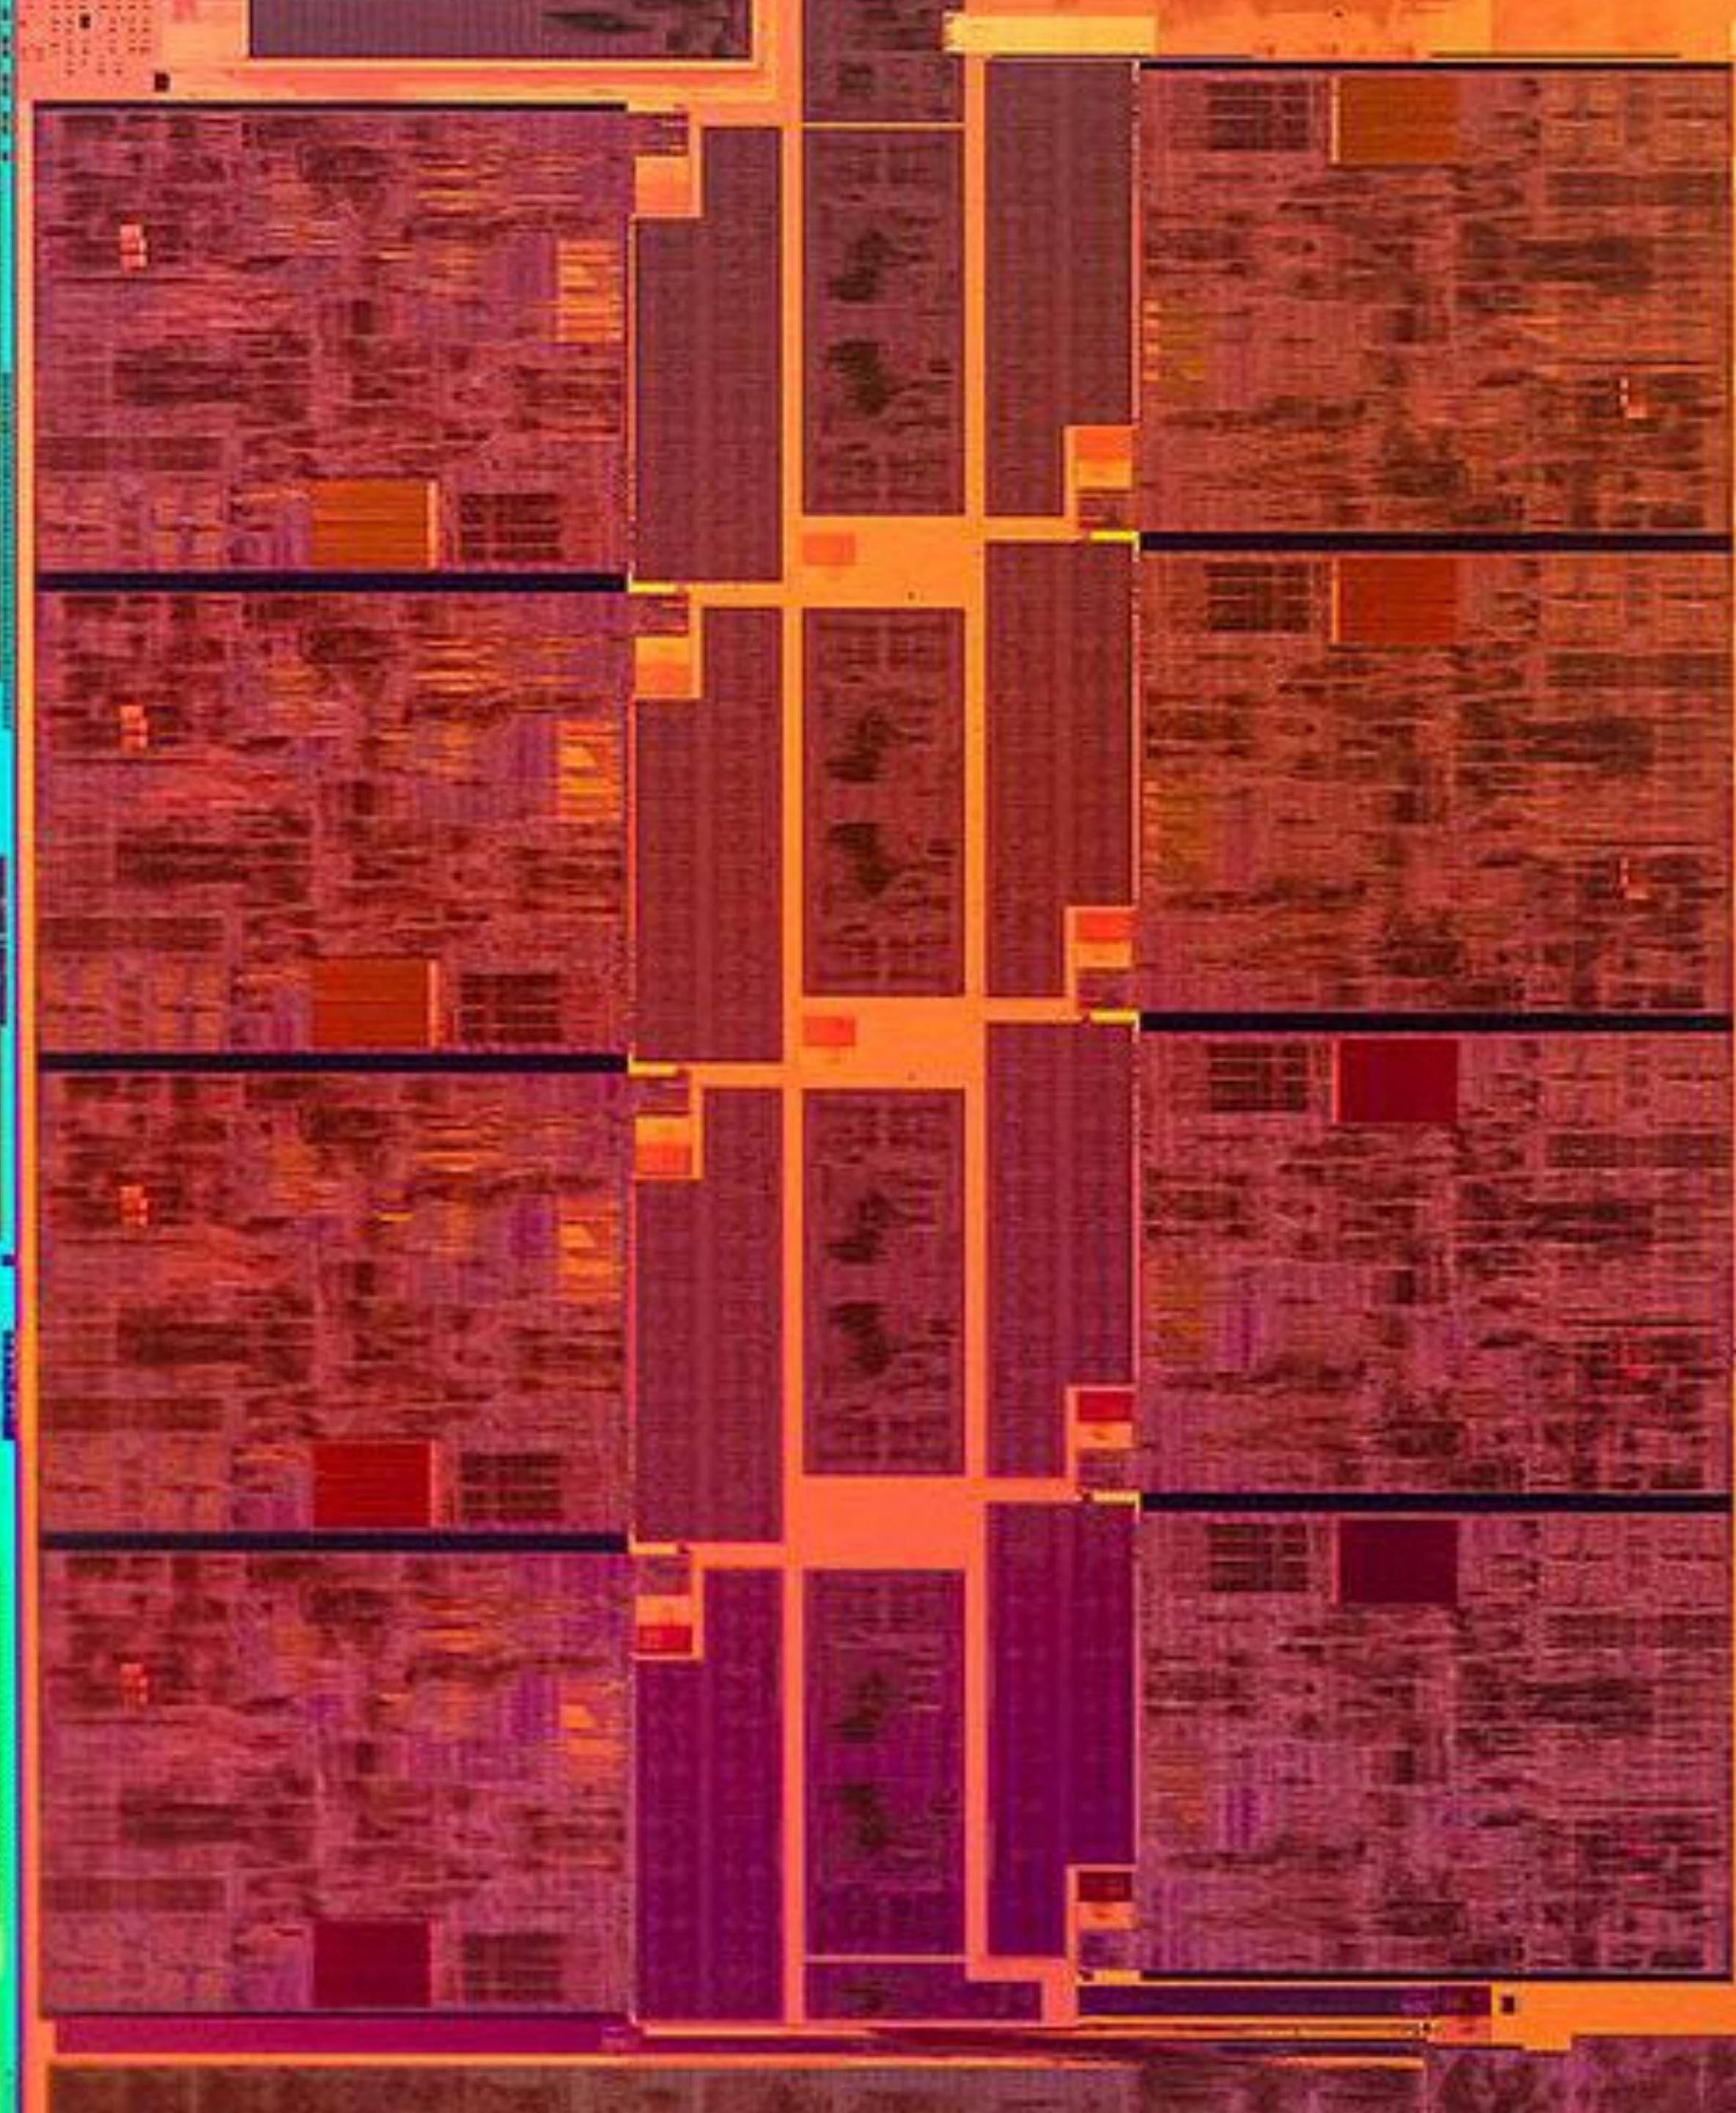
“Brain Float”

- 1 data type:
 - 32* 16 bit truncated FP32s, AKA “BF16”
 - AI firms using HPC loved this one for training workloads

No fun pictures, Cooper lake had a limited installation market

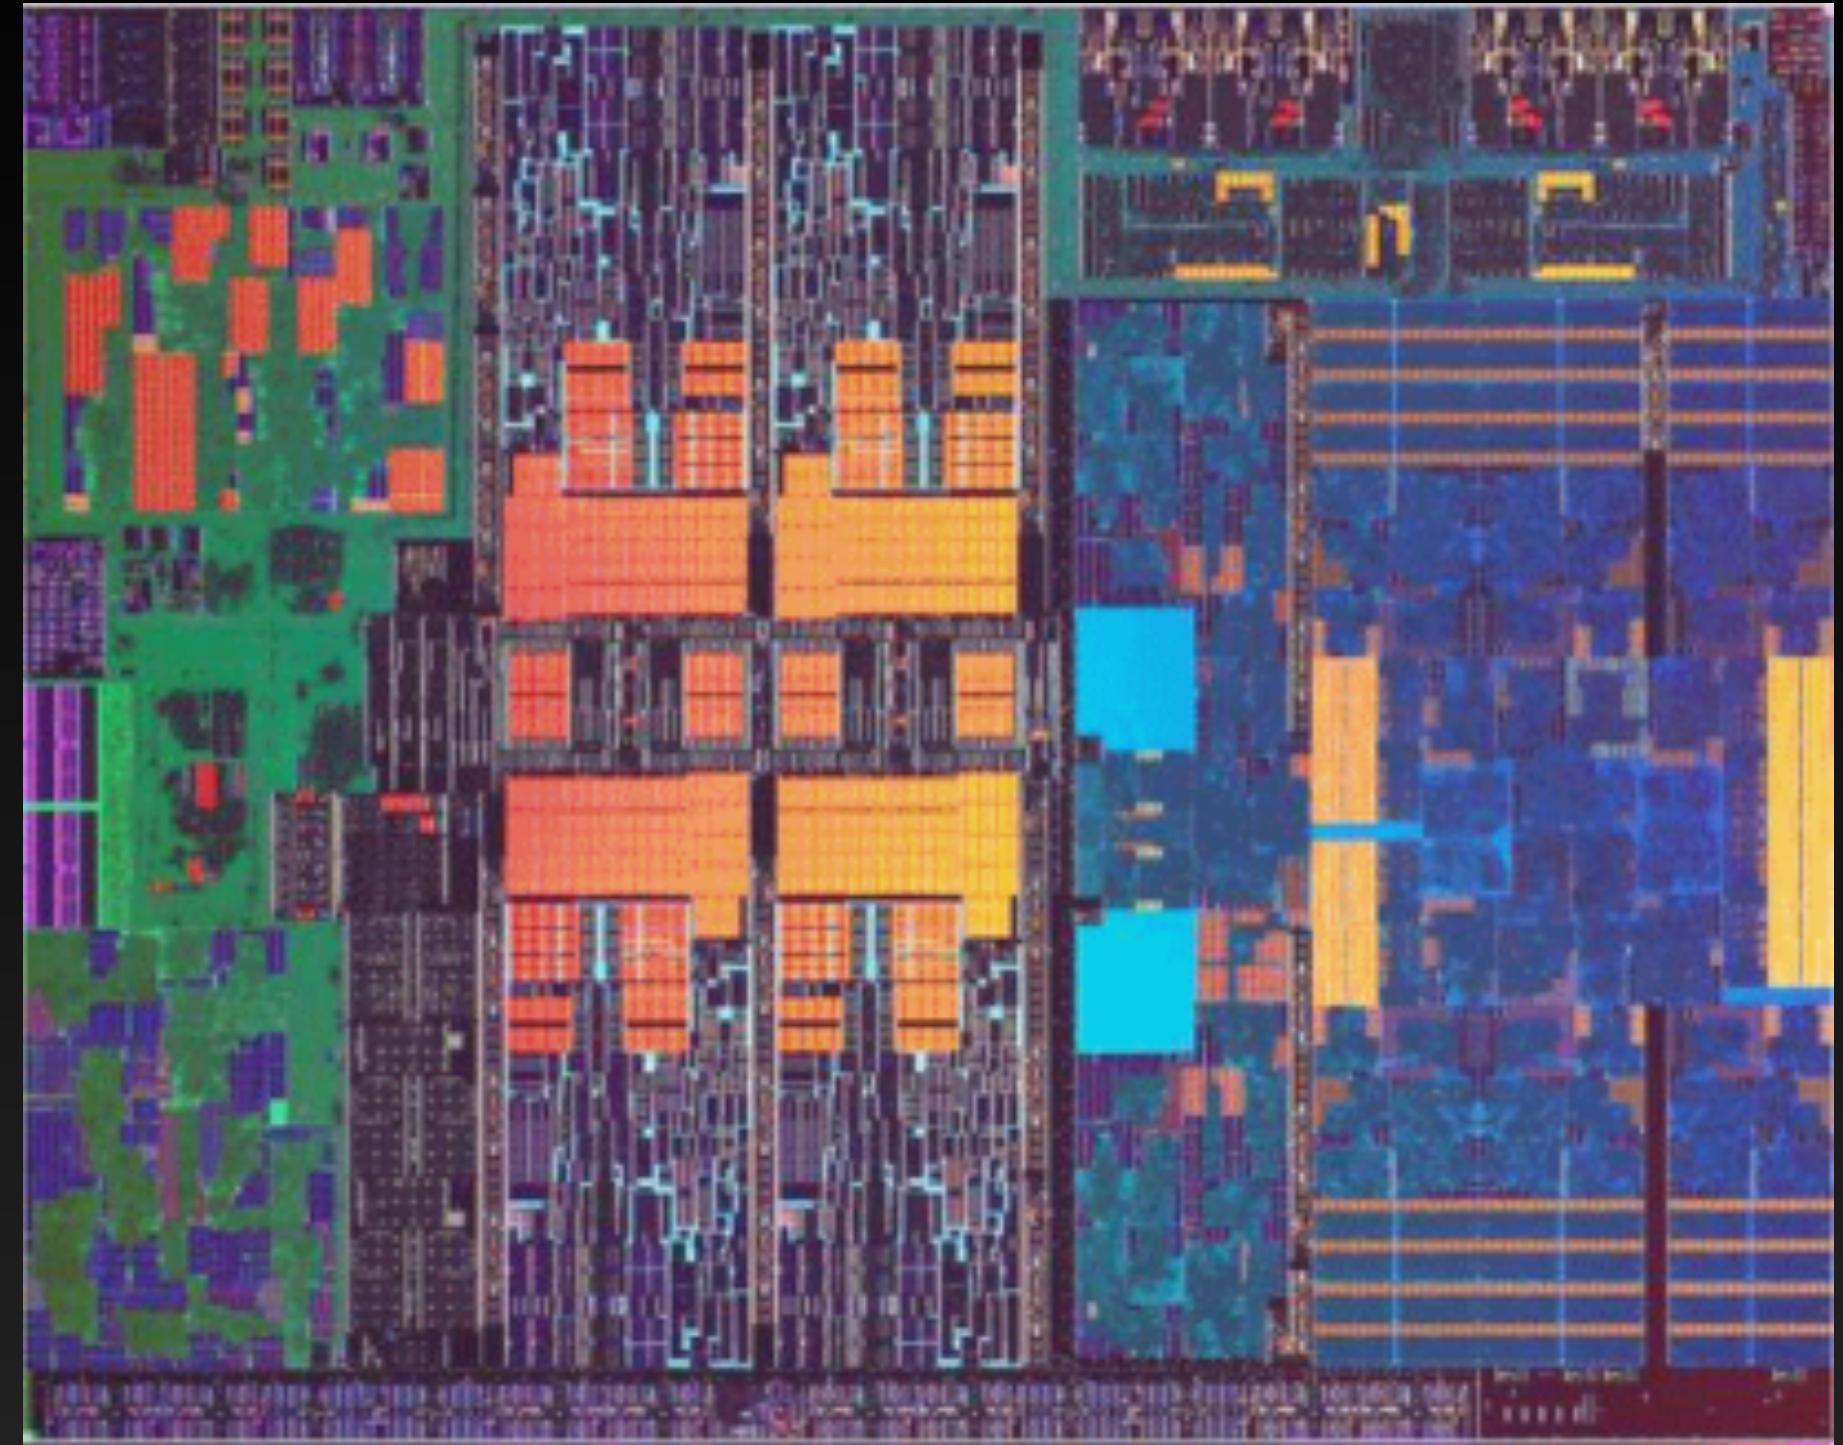
AVX512 “V7” “VNNI, VBMI2, BITALG”

- 0 data types
- Basically integer FMA for (u)int8 and (u)int16
- AI using HPC loved this one for CPU inference



AVX512 “V8” VP2INTERSECT

- 0 data types
- Micro coded Instructions to compute the intersection of 2 vectors and return a mask
- Devs proved post release it could be done faster with “basic” AVX512 instructions (bitalg, VBMI etc.)
- Has since been deprecated



 @fclc@mast.hpc.social @FelixCLC_ · Jul 19
Hey @IntelDevTools and @IntelDevTools!

Trying to confirm if we should consider the #AVX512 VP2intersect extensions as deprecated in current and future chips for SW roadmaps. Any chance I can get a confirmation?

It's been removed from future chip features in docs 😊😊

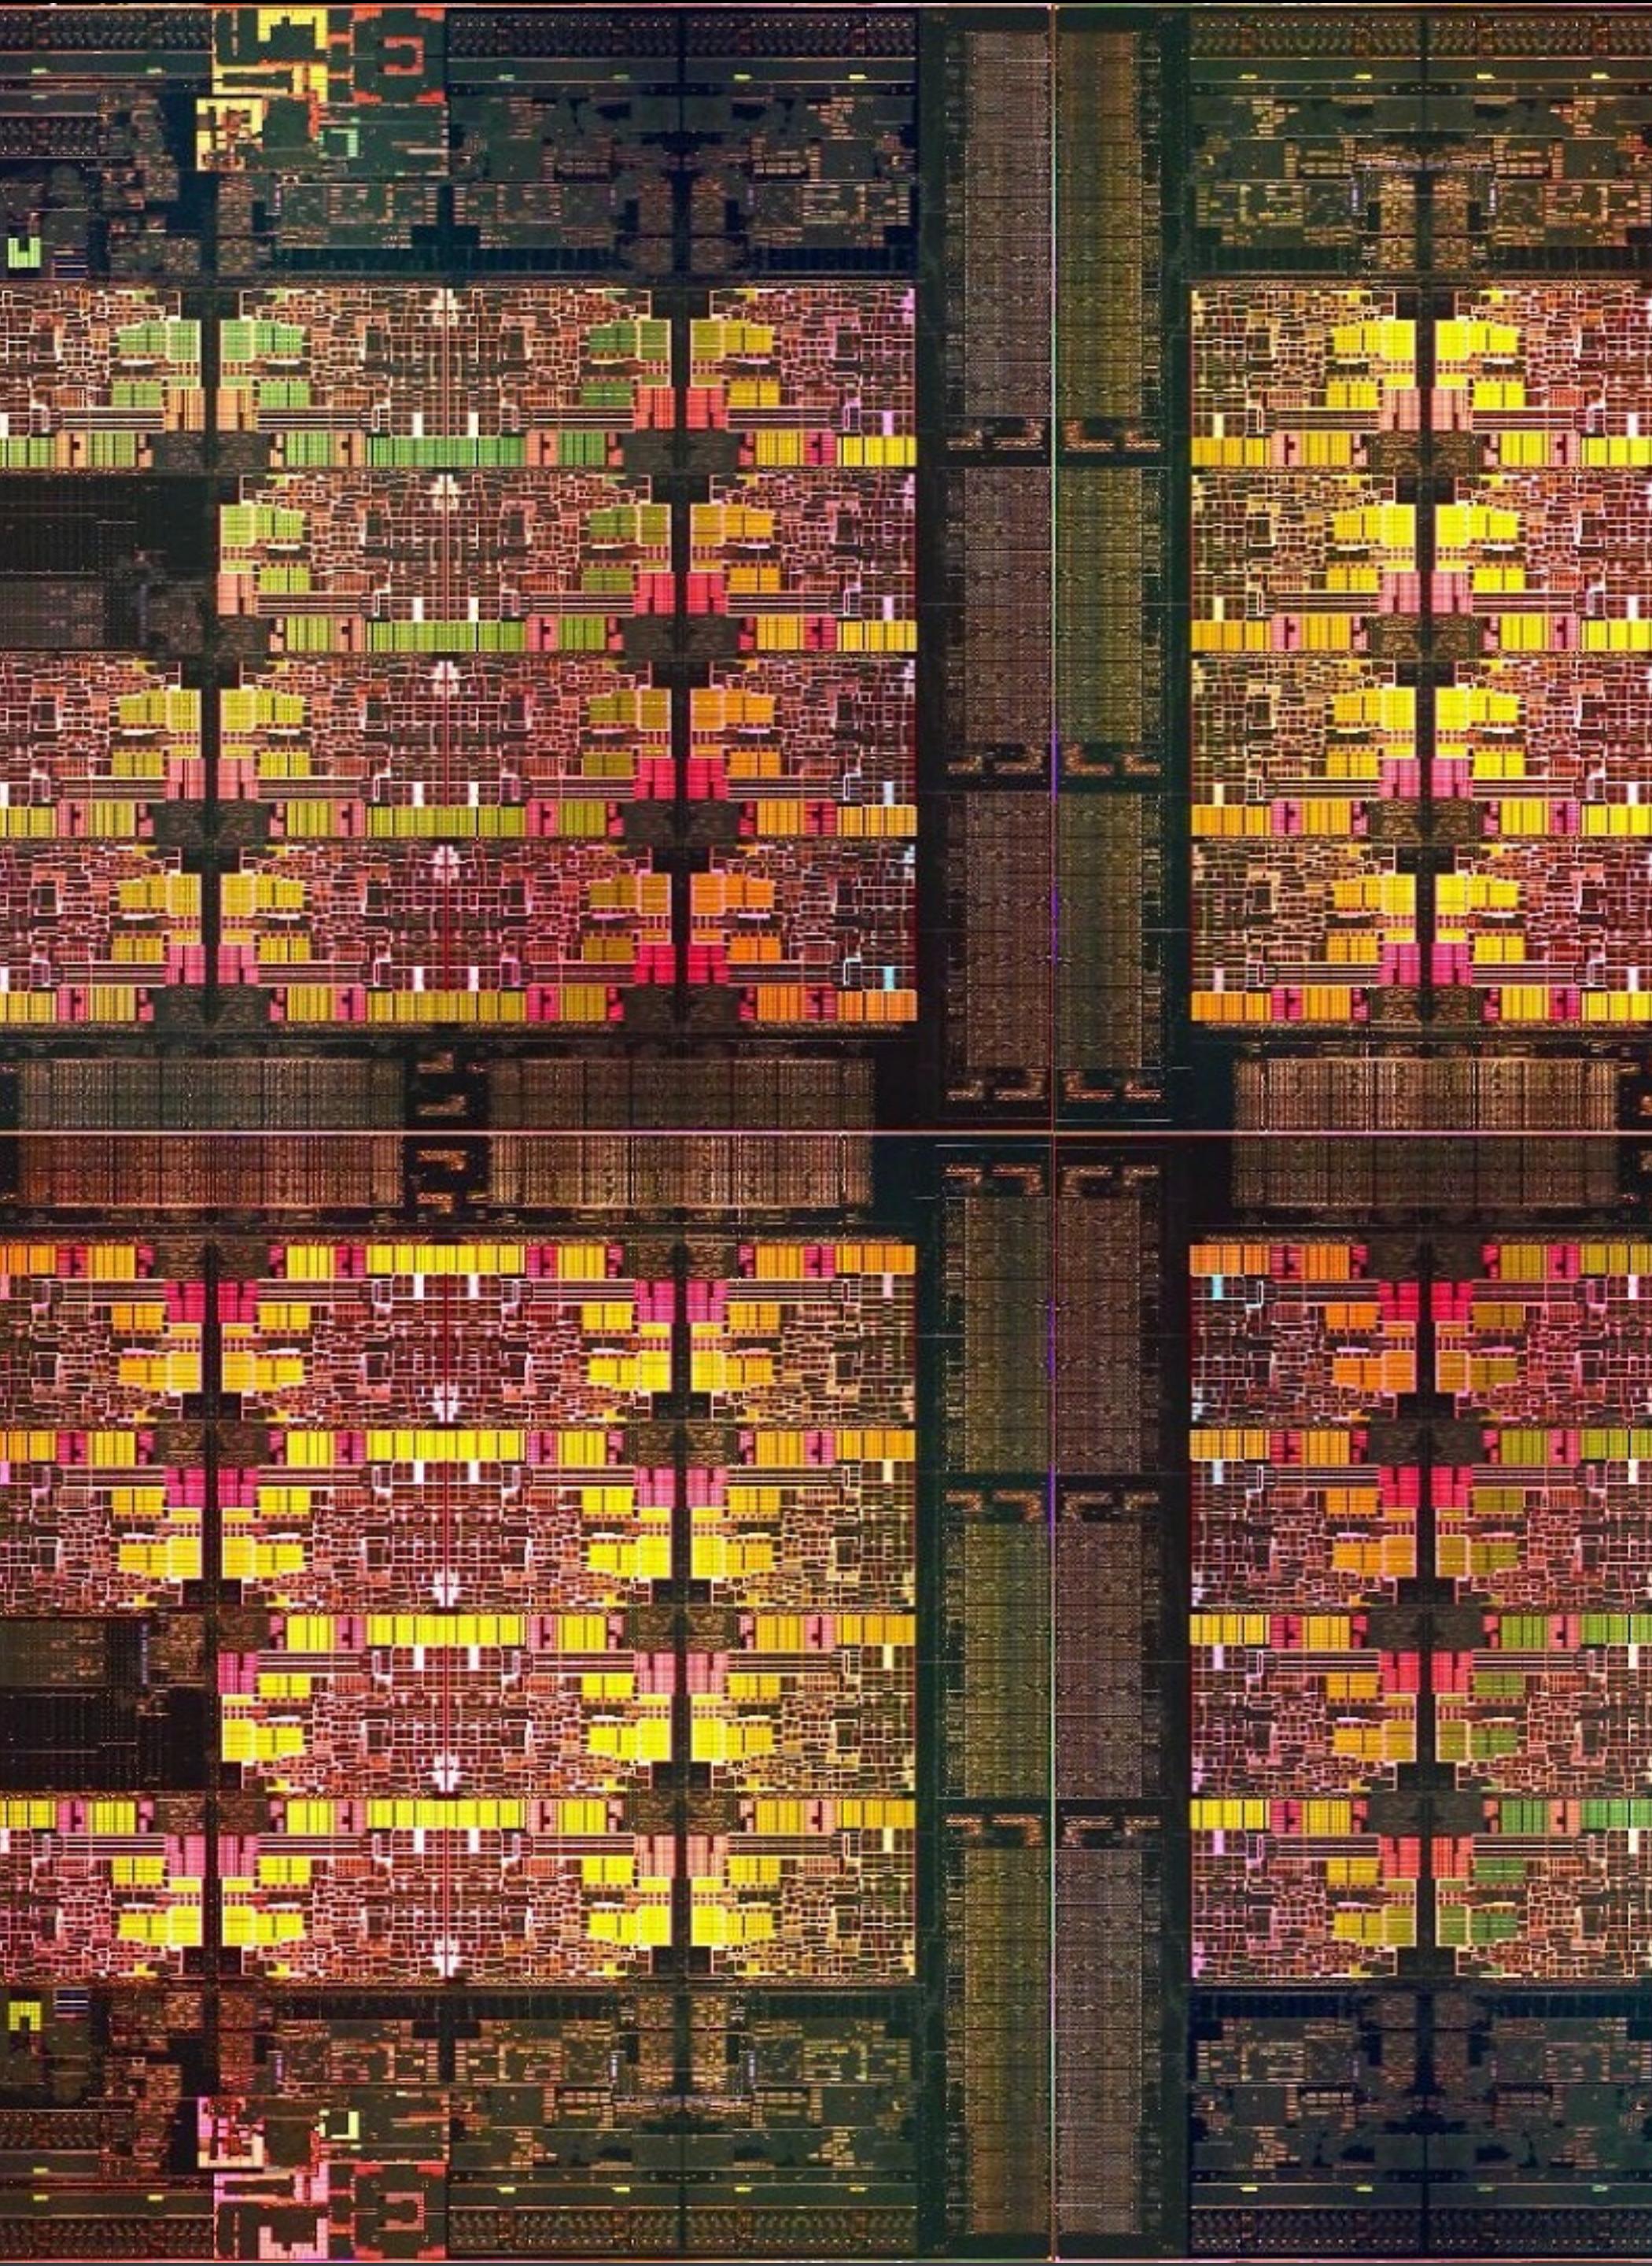
3 1 7 1,911 ⌂ ...

 Ronak Singhal @rsinghal1

Yes. Deprecated.

AVX512 “V9” FP16

- 1 + i data types
- IEEE754-2008 compliant Binary16
- Provides everything that FP32 and FP64 have
- Even gets its own series of complex number Instructions
- HPC: Good/great for mixed precision algorithms, includes FP16 compute on top of storage
- Expands F16C to also be able to move from FP16<->FP64 directly, instead of store f16C:
Load FP16->FP32->FP64-> math -> FP32->FP16 store



**Targets that implement all the way to
AVX512{F,BW,CD,DQ,VL} & are primary
processors are considered x86-64-V4
compliant**

All Intel HPC CPU since ~2015 (SKX)
AMD HPC Processors since ~2021 (Zen4)
For the 2 of you in the room, Centaur CNS (doesn't matter)

Really powerful data processing, lots of massive vectors

So AVX10?

Lots of talk about what AVX512 was, what in the world is AVX10?

Section 5

No more **subversioning**

expectation

The goal of AVX10 is to put an end to this:

"I hope AVX512 dies a painful death, and that Intel starts fixing real problems instead of trying to create magic instructions to then create benchmarks that they can look good on.

[...]

Stop with the special-case garbage, and make all the core common stuff that everybody cares about run as well as you humanly can. Then do a FPU that is barely good enough on the side, and people will be happy. AVX2 is much more than enough.

Yeah, I'm grumpy."

Linus Torvalds

And this:

SIMD: 512-bit

* AVX512-F	= Yes
AVX512-CD	= Yes
AVX512-PF	= No
AVX512-ER	= No
* AVX512-VL	= Yes
* AVX512-BW	= Yes
* AVX512-DQ	= Yes
AVX512-VPOPCNTDQ	= Yes
AVX512-4VNNIW	= No
AVX512-4FMAPS	= No
* AVX512-IFMA	= Yes
* AVX512-VBMI	= Yes
AVX512-VNNI	= Yes
* AVX512-VBMI2	= Yes
AVX512-BITALG	= Yes
AVX512-VPCLMULQDQ	= Yes
AVX512-VAES	= Yes
* AVX512-GFNI	= Yes
AVX512-BF16	= Yes
AVX512-VP2INTERSECT	= Yes
AVX512-FP16	= Yes

And this:

```
global GetAVX512Capability
global GetAVXCapability
global GetSSECapability
global GetAVX512fp16
global Getfp16c

section .text

; Checks CPUID for AVX capability
GetAVXCapability:
    push rbx
    mov eax, 1
    cpuid
    mov ecx, eax
    shr eax, 28          ; Read bit 28
    and eax, 1
    pop rbx
    ret

; Checks CPUID for SSE capability
GetSSECapability:
    push rbx
    mov eax, 1
    cpuid
    mov edx, eax
    shr eax, 25          ; Read bit 25
    and eax, 1
    pop rbx
    ret

; Checks CPUID for AVX512 capability
GetAVX512Capability:
    push rbx
    mov eax, 7
    xor ecx, ecx
    cpuid
    mov eax, ebx
    shr eax, 16           ; Read bit 16
    and eax, 1
    pop rbx
    ret

; Checks CPUID for AVX512-fp16 capability
GetAVX512fp16:
    push rbx
    mov eax, 7
    xor ecx, ecx
    cpuid
    mov edx, eax
    shr eax, 23           ; Read bit 23 of register edx
    and eax, 1
    pop rbx
    ret

Getfp16c:
    push rbx
    mov eax, 1
    xor ecx, ecx
    cpuid
    mov ecx, eax
    shr eax, 29           ; Read bit 29 of register ecx
    and eax, 1
    pop rbx
    ret
```

And this:

```

case 5: // Skylake X
    if(support_avx512_bf16())
        return CPUTYPE_COOPERSLAKE;
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 14: // Skylake
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 7: // Xeon Phi Knights Landing
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 12: // Apollo Lake
case 15: // Denverton
    return CPUTYPE_NEHALEM;
}
break;
case 6: // family 6 exmodel 6
switch (model) {
case 6: // Cannon Lake
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
}

```

```

if(support_avx())
    return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 10: // Ice Lake SP
    if(support_avx512_bf16())
        return CPUTYPE_COOPERSLAKE;
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 7: // family 6 exmodel 7
switch (model) {
case 10: // Goldmont Plus
    return CPUTYPE_NEHALEM;
case 14: // Ice Lake
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 8:
switch (model) {
case 12: // Tiger Lake
case 13: // Tiger Lake (11th Gen
Intel(R) Core(TM) i7-11800H @ 2.30GHz)
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())

```

```

        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 14: // Kaby Lake and refreshes
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 15: // Sapphire Rapids
    if(support_avx512_bf16())
        return CPUTYPE_COOPERSLAKE;
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 10: //family 6 exmodel 10
switch (model) {
case 5: // Comet Lake H and S
case 6: // Comet Lake U
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 9:
switch (model) {
case 7: // Alder Lake desktop
case 10: // Alder Lake mobile
    if(support_avx512_bf16())
        return CPUTYPE_COOPERSLAKE;
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 7: // Rocket Lake
if(support_avx512())
    return CPUTYPE_SKYLAKEX;
if(support_avx2())
    return CPUTYPE_HASWELL;
if(support_avx())
    return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
}
```

```

if(support_avx())
    return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
case 14: // Kaby Lake and refreshes
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 10: //family 6 exmodel 10
switch (model) {
case 5: // Comet Lake H and S
case 6: // Comet Lake U
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 9:
switch (model) {
case 7: // Alder Lake desktop
case 10: // Alder Lake mobile
    if(support_avx512_bf16())
        return CPUTYPE_COOPERSLAKE;
    if(support_avx512())
        return CPUTYPE_SKYLAKEX;
    if(support_avx2())
        return CPUTYPE_HASWELL;
    if(support_avx())
        return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
case 7: // Rocket Lake
if(support_avx512())
    return CPUTYPE_SKYLAKEX;
if(support_avx2())
    return CPUTYPE_HASWELL;
if(support_avx())
    return CPUTYPE_SANDYBRIDGE;
else
    return CPUTYPE_NEHALEM;
}
break;
}
```

And this:

The screenshot shows a GitHub repository page for 'vpu-count'. The repository is public and has 4 watchers, 3 forks, and 33 stars. It features a master branch with 2 branches and 1 tag. The repository is described as providing information about AVX-512 support on recent Intel processors. It includes files like .gitignore, Makefile, README.md, empirical.c, intel-xeon-scalable-sku.txt, test.c, time.c, and vpu-count.c. The last commit was made by jeffhammond on April 10, 2022, with 63 commits.

vpu-count Public

Watch 4 Fork 3 Star 33

master 2 branches 1 tag

Go to file Add file Code

jeffhammond fix ice lake server and tiger lake detection; add sapphire rapids 24d11ba on Apr 10, 2022 63 commits

File	Description	Time Ago
.gitignore	Icelake client (#10)	4 years ago
Makefile	fix ice lake server and tiger lake detection; add sapphire rapids	last year
README.md	add Ice Lake Xeon data now that it is public	2 years ago
empirical.c	update empirical code (#9)	4 years ago
intel-xeon-scalable-sku.txt	improve things	5 years ago
test.c	improve things	5 years ago
time.c	improve things	5 years ago
vpu-count.c	fix ice lake server and tiger lake detection; add sapphire rapids	last year

About

Information about AVX-512 support on recent Intel processors

xeon floating-point avx512
avx-instructions fused-multiply-add
intel-processor

Readme Activity 33 stars 4 watching 3 forks Report repository

I think everybody gets
the point

I think everybody gets
the point

If not, don't know what to tell you

So AVX10?

Answer the darn question!

AVX10 is actually AVX10.N/M

What does the .N in AVX10.N/M mean?

Simple: N is version number. Something get added? N increments

Importantly: if you have N+4, you ***must*** support N+3, N+2 ... N

What does the /M in AVX10.N/M mean?

Simple: M is *implementation size*

Valid Options are 512, 256, 128*

Building a small consumer CPU? Make AVX10.N/
256. Smaller registers, less state, less power.

Importantly: must support all smaller sizes

What does AVX10.N/M include?

Version 1, AKA AVX10.1/M supports:

All of AVX2, AVX1, FMA, SSE1-4.2

and

All of the **instructions** from AVX512 {F, CD, VL, DQ, BW,
IFMA, VBMI, VBMI2, VPOPCNTDQ, BITALG, VNNI,
VPCLMULQDQ, GFNI, VAES, BF16, FP16}

The italics and asterix on that last slide are
very  , what's up with that?

Because AVX10 is meant to allow instruction
support on smaller chips, it keeps the
instructions from AVX512, but not the register
size

Isn't that just going to make things worse?

Yes and no

It means you can run a single check and

know what is and isn't supported

Does mean you need to check local

implementation size

Say again?

Most of the Instructions from {AVX512} are great for commercial data centre as well as local client data processing; they don't do much for classical HPC FP64 apps

So AVX10 for HPC?

Beyond the spec, WHAT IS AVX10?!?!

**For admins and users:
A single compiler flag to enable all of the
goodies in the history of AVX512**

For programmers/developers/RSE's:
A pair of CPUID flags to check for kernel
version targeting and implementation size

Is it out?

Yes* 😊

Section 6

AVX10 for HPC in 2023

**AVX10.1/512 is coming next year in silicon
with Intels Granite Rapids Xeons**

DON'T CARE HOW

AV

con

I WANT IT NOW

**GOOD NEWS
EVERYONE!**



**AVX10.1/512 is the same as Sapphire
Rapids and upcoming Emerald
Rapids**

**For those with intel dev cloud or ANL/
ALCF Sunspot access, your codes for
SPR are the same ISA as AVX10.1/512**

Section 7

That's all folks!

Brief note on AVX10/256:

**AVX10/256 devices are coming soon.
Personally I expect Intel in '24, AMD in '25**

Brief note on AVX10/128

AVX10/128 is a silly, silly idea that should never have been allowed to exist, and that I've asked intel to remove from the spec.

Brief note on AVX10/128

I wrote an entire companion article on this topic for the publication Chips and Cheese

Link: <https://chipsandcheese.com/2023/10/11/avx10-128-is-a-silly-idea-and-should-be-completely-removed-from-the-specification/>



Information:

Mostly active on mastodon

fclc@mast.hpc.social

Twitter (for now) @felixclc_

Email:

felix.leclair123@hotmail.com



Acknowledgements:

Everyone who provided quotes

Various CPU architects who provided guidance

Various “spec lawyers” who provided insights

Chips and Cheese for hosting the accompanying

article

Acknowledgements: Easy Build

**Now the fun Part:
Closed session Q and A**